

>>> network .toCode()

# Workshop Telemetry Deepdive for Baxter

*Nov 19th and Nov 20th*

# Agenda Day 1

- Introduction to Telemetry and Metrics
- Target Architecture + Lab
- PromQL & Metrics
  - Prometheus and PromQL Overview
  - Queries and Filtering
  - Functions
  - Operators

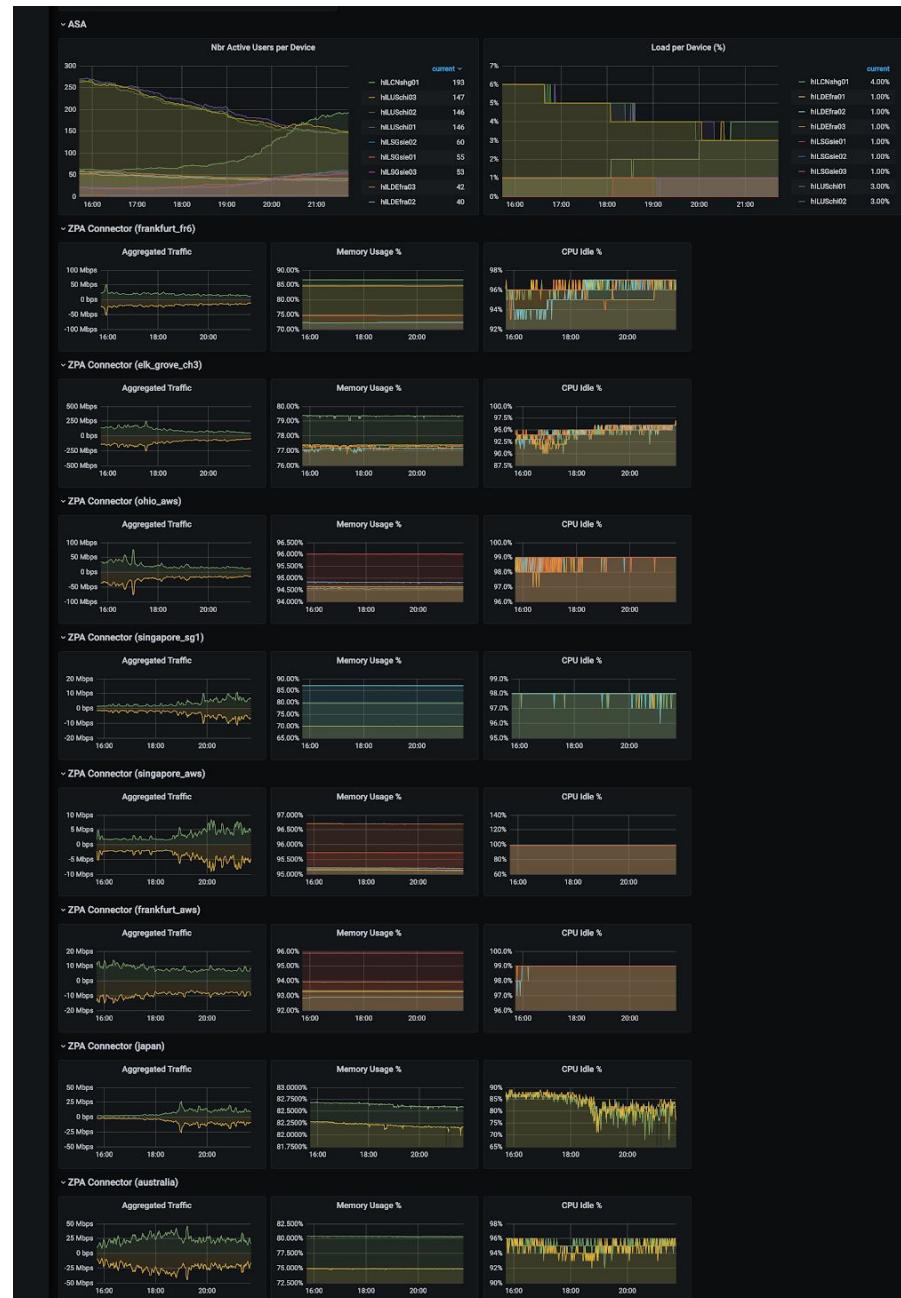
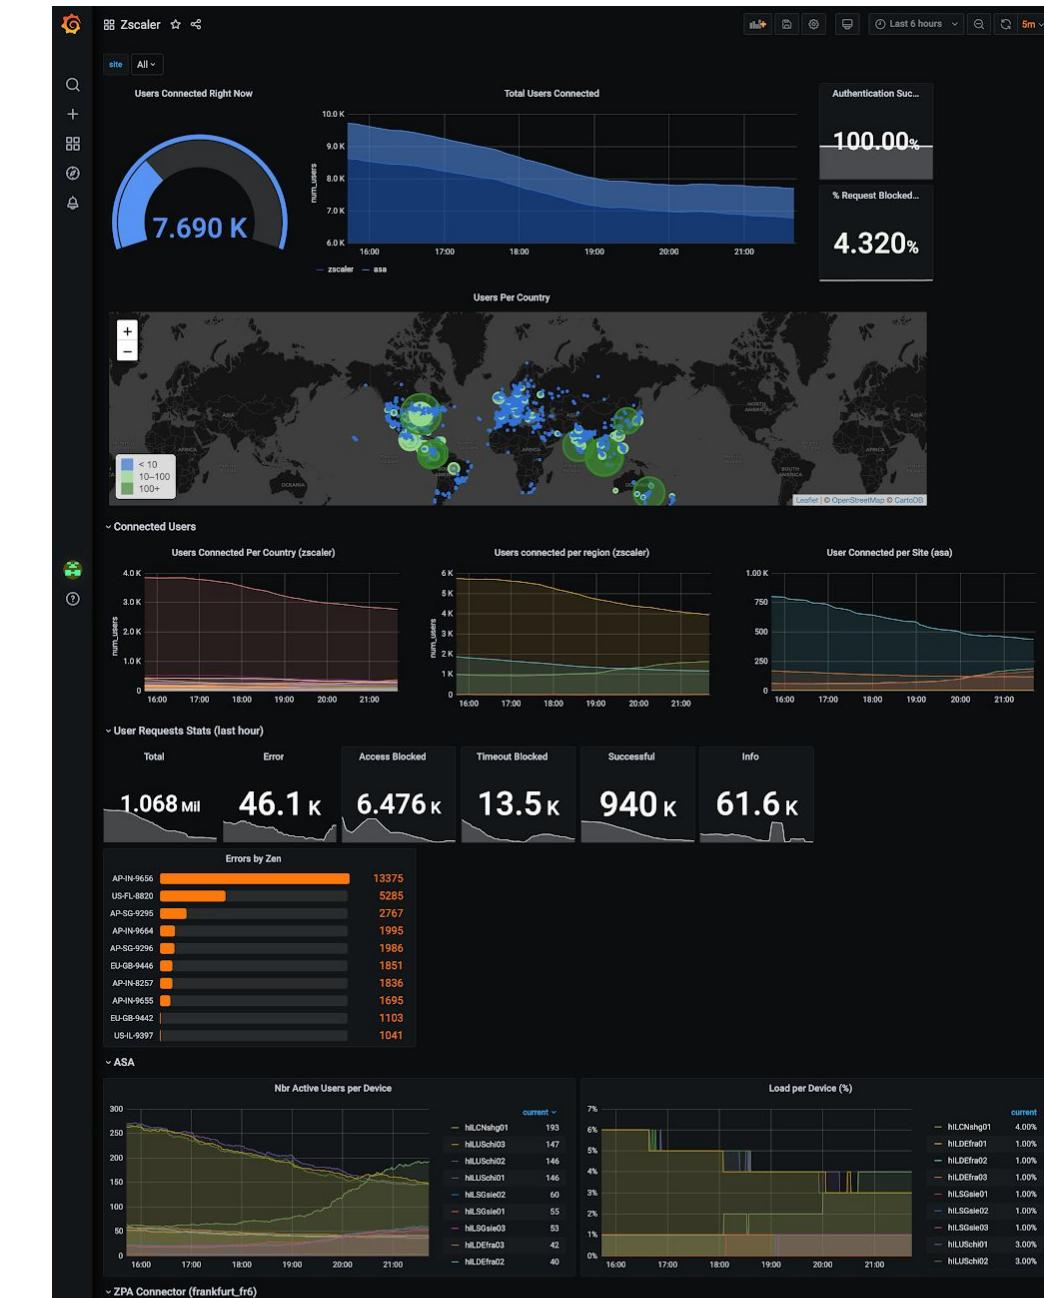
# Agenda Day 2

- Recap from Day 1
- Introduction to Grafana
- Variables
- Additional Panels
- Annotations
- Alertings

```
>>> network .toCode()
```

# Day 1

# Introduction to Telemetry and Time Series Database



# Myths around Network Monitoring

**#1 - My network monitoring sucks because of SNMP**

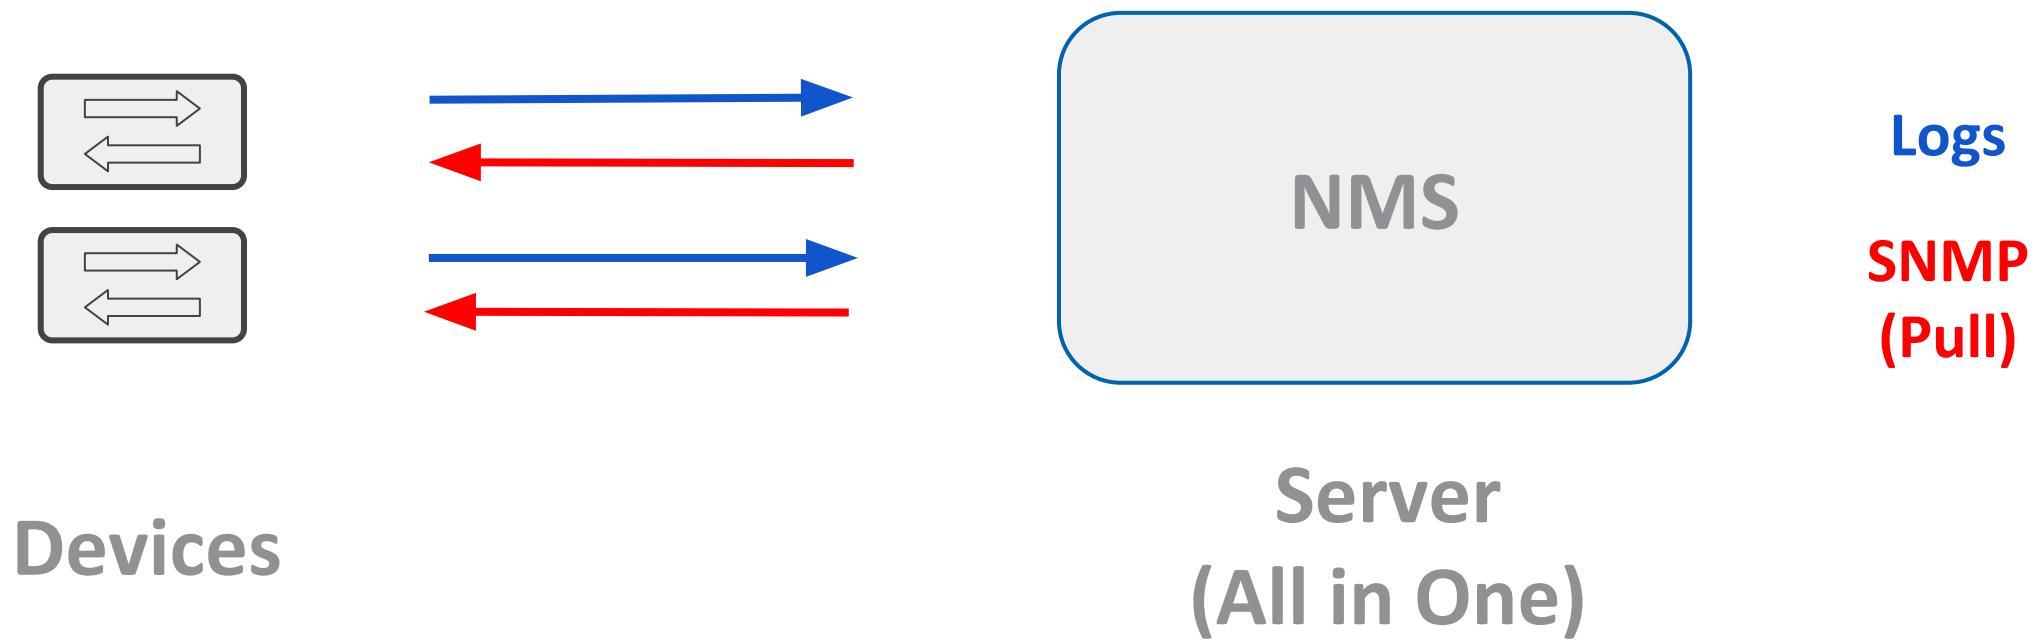
**#2 - Telemetry Streaming is the solution to my problem**

**#3 - Faster is better**

**#4 – I need to get rid of SNMP**

# **Myth #1 : My network monitoring sucks because of SNMP**

# Legacy Network Monitoring Solution



What tools are you using ?

OpenNMS  
Collectd

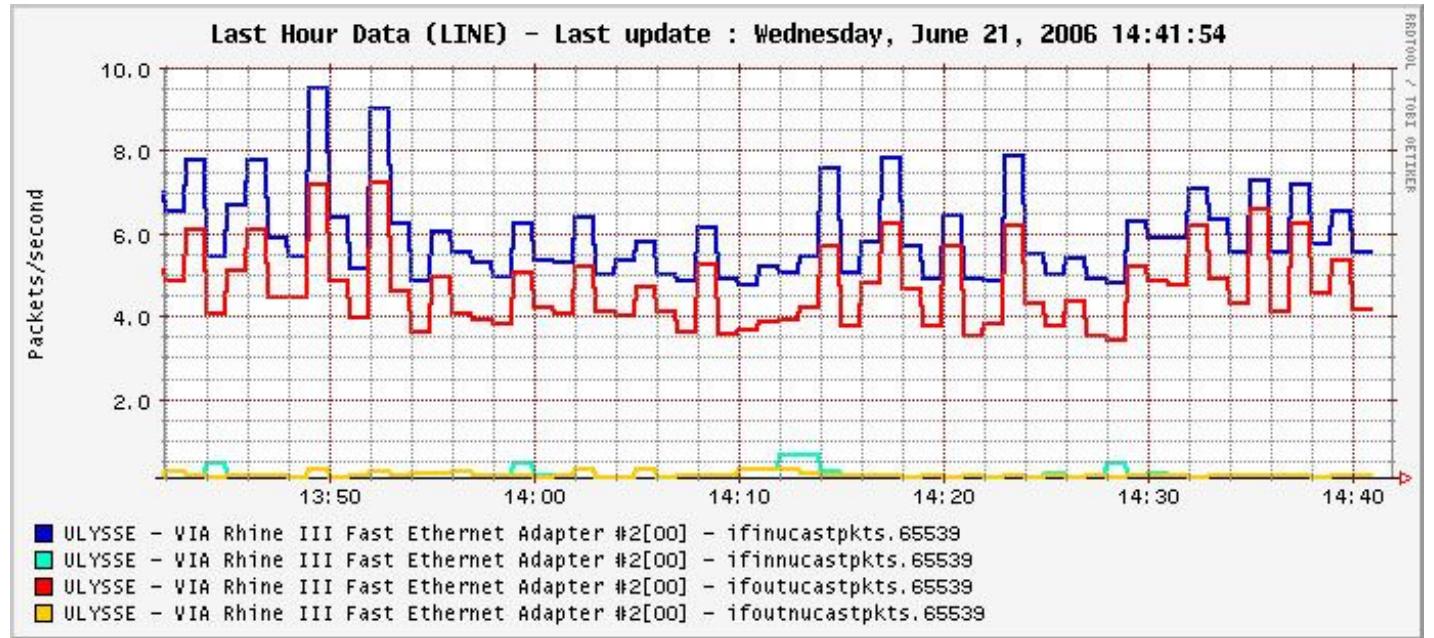
**RRDtools**

ZObservium  
Cacti  
MRTG  
Nagios  
Solaris

# RRD Tools

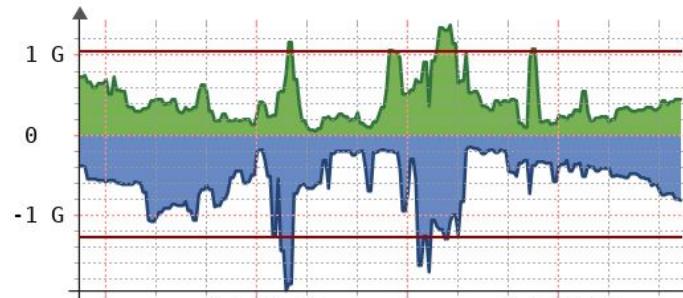
- Introduced in 1999
- Storage
- Aggregation
- Visualization

No query engine  
Data retention is poor.

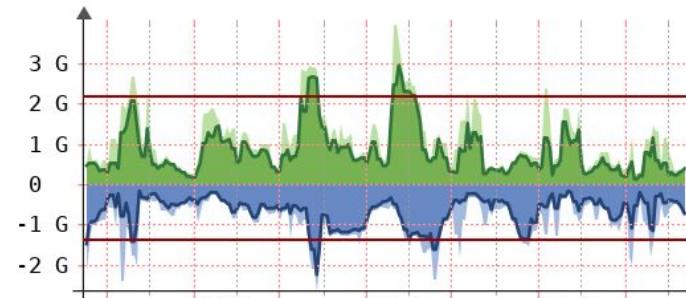


# RRD Tools - Down Sampling

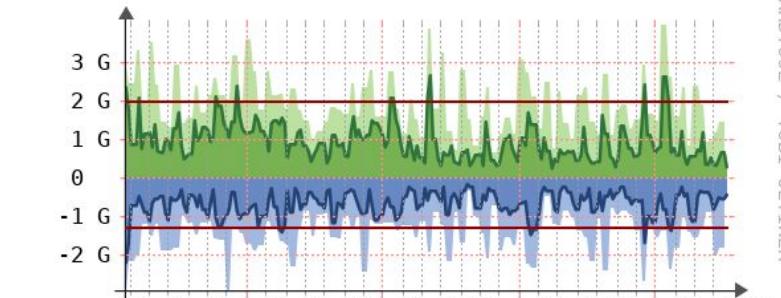
## Traffic



RRDTOOL / TOBI OETIKER



RRDTOOL / TOBI OETIKER



RRDTOOL / TOBI OETIKER

1 Day

1 Week

1 Month

# **Myth #2 : Telemetry Streaming is the solution to my problem**

**Telemetry has been a hot topic in the network industry**

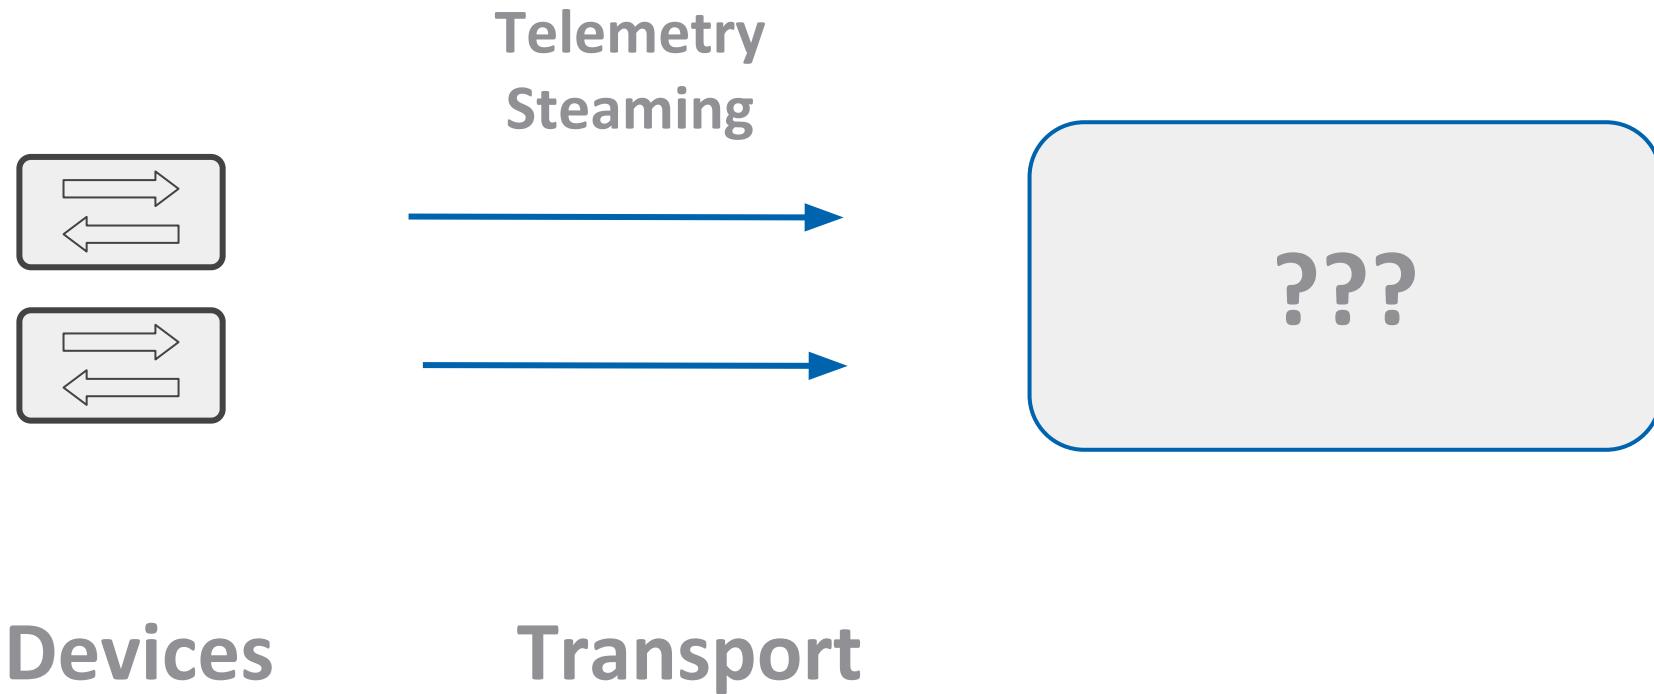
Telemetry  
Streaming

Kill SNMP

gNMI

Openconfig

# ... Network Monitoring Solution



PULL

PUSH



SNMP

Streaming

# History of Telemetry Streaming

## 1<sup>st</sup> Generation

Traffic send from the PFE directly

Google Protocol Buffer over UDP

Very fast  
Very Efficient

## 2<sup>nd</sup> Generation

Traffic send from the RE

Google Protocol Buffer over gRPC

Proprietary interface

## 3<sup>rd</sup> Generation

Traffic send from the RE

Google Protocol Buffer over gRPC (GPB-KV)

gNMI interface (standard)

# Use Cases for gNMI / Telemetry Streaming

- SDN Controller (centralized Management and Control plane)
- High frequency monitoring (less than 30s)
- Event notification (Interface Up/Down, BGP etc ..)
- Extract State

# Myth #3 : Faster is better

# Type of data in a monitoring

## Metrics

..

## Time Series

Numeric value evolving  
over time

Constant Interval

Counters

CPU

Number peers

## Logs

## Events

Mostly Text data

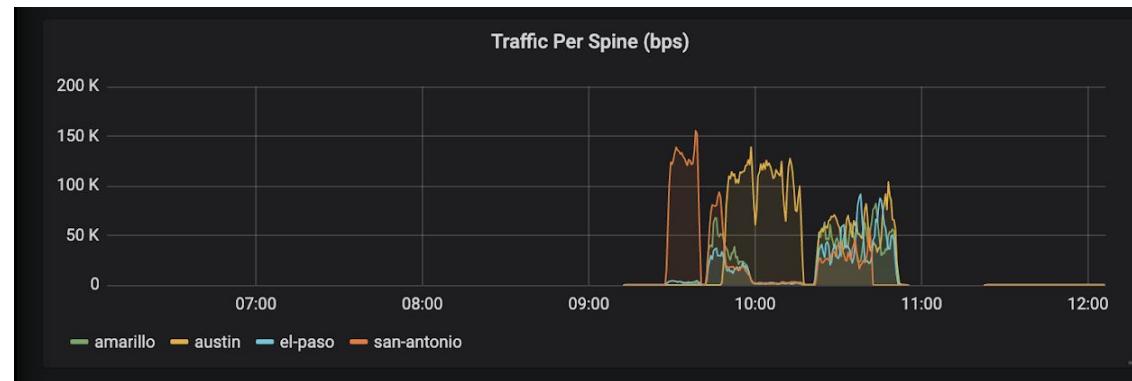
Unpredictable interval

## Structured Data

Routing/Forwarding Table  
Configuration

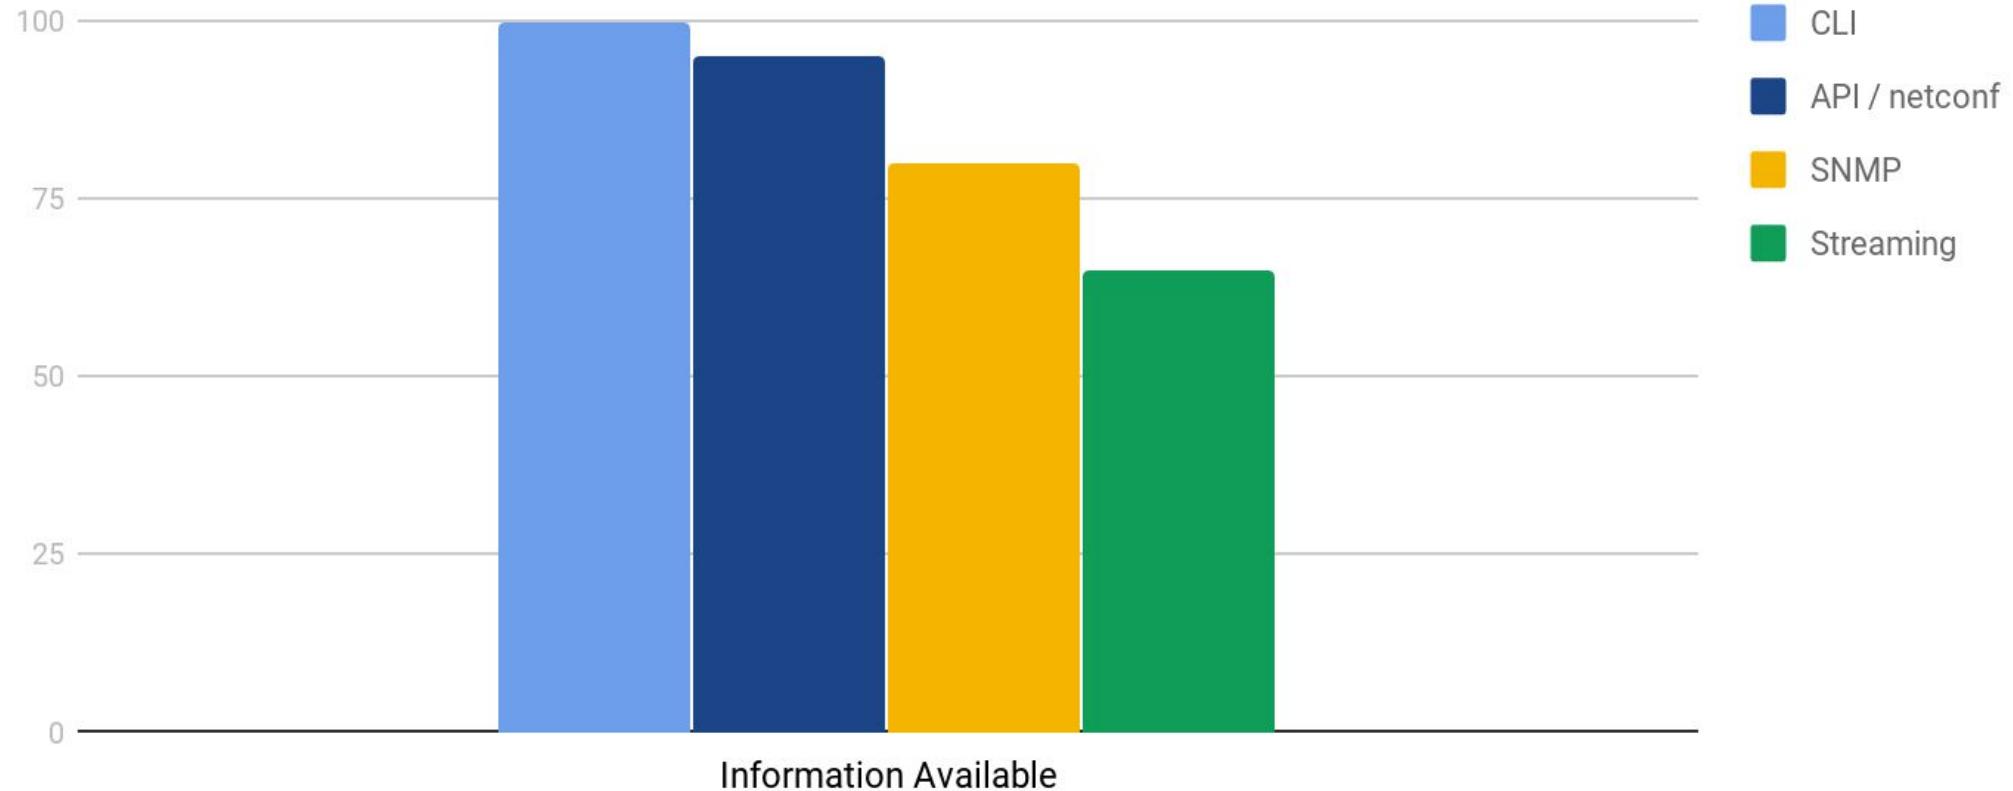
# Metrics data needs consistency more than speed

- Time series database expect consistent datapoints.
- Most databases won't be able to ingest counters every 5s
- Most network devices won't be able to generate data every 5s
- Most network engineers won't be able to consume update that frequently

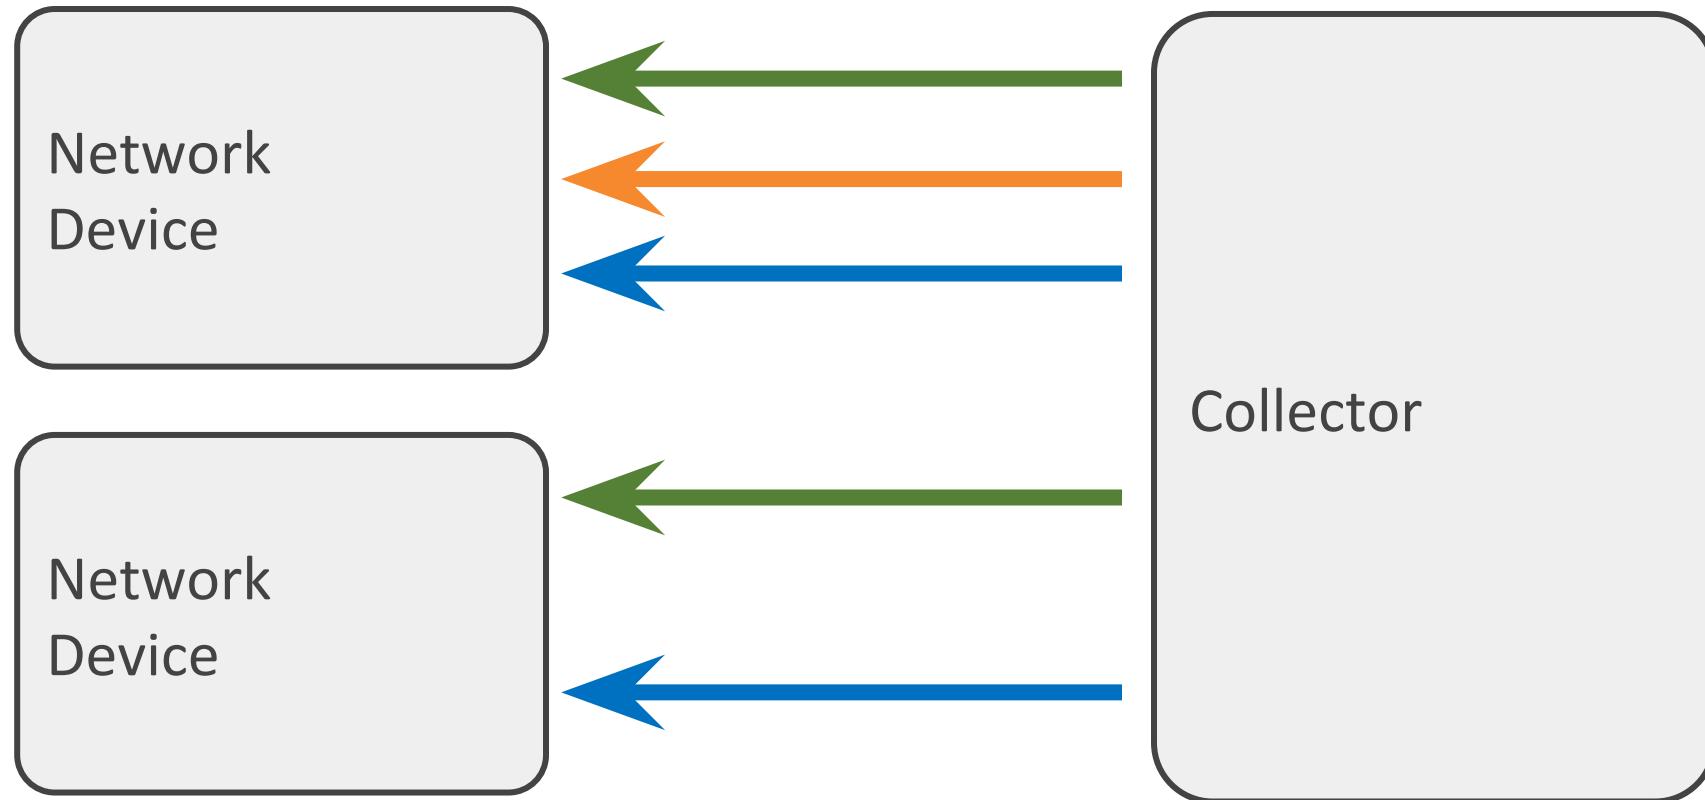


# **Myth #4 : I need to get rid of SNMP**

# Volume of Data available



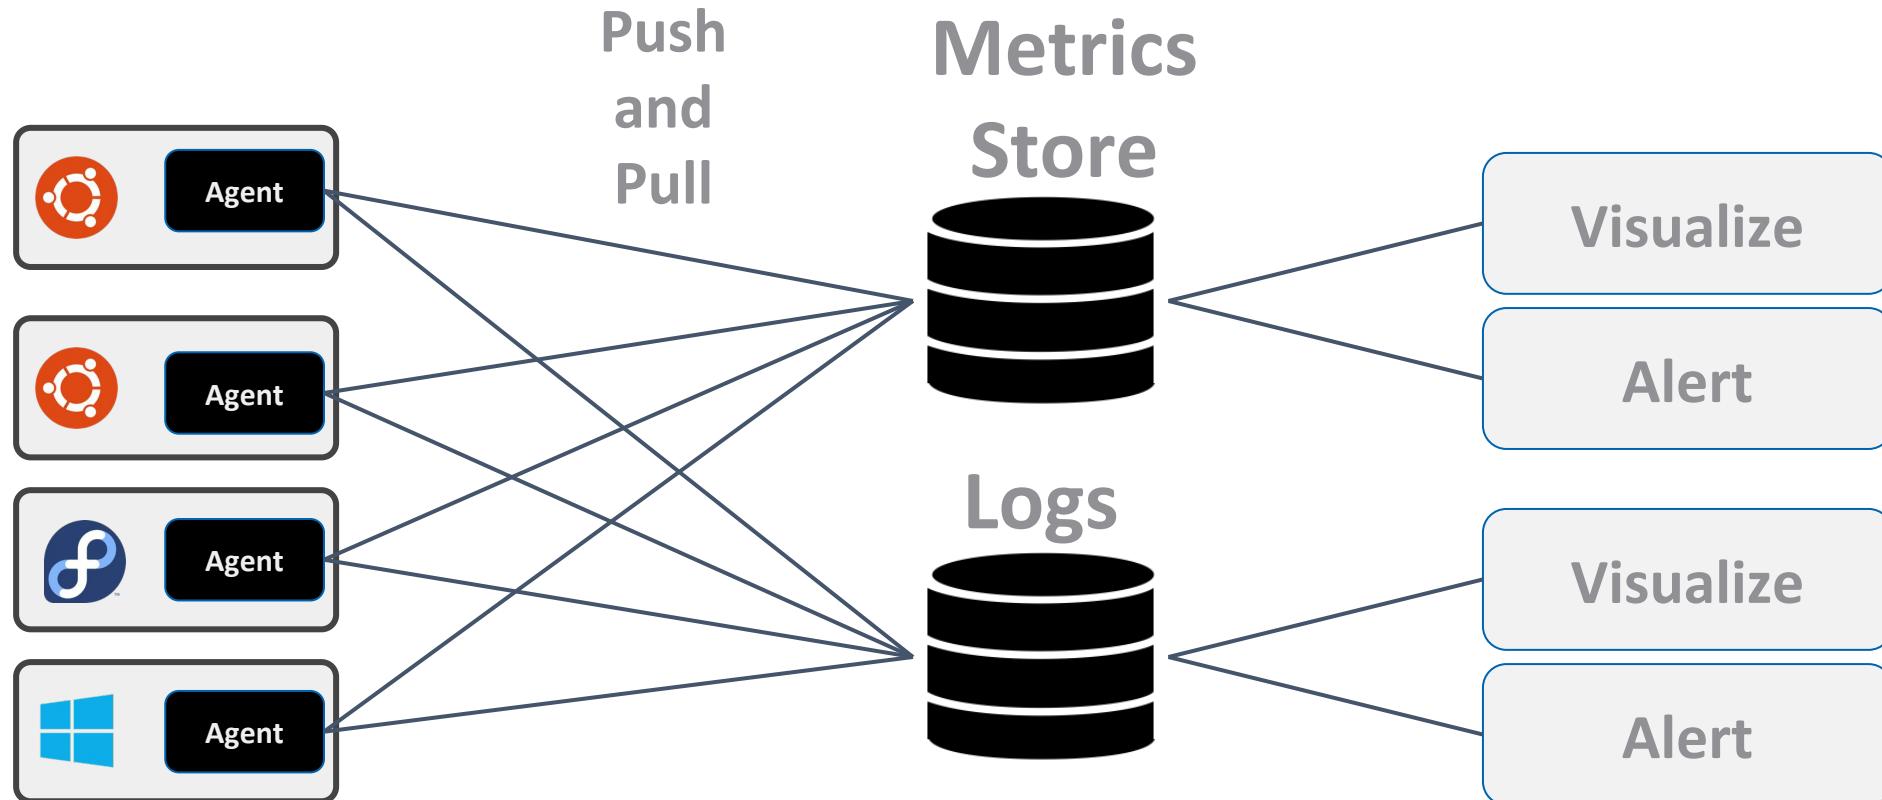
# Embrace all of them



2

## Monitoring outside of the network industry

# What are others doing outside of the Network Industry ??



- Each component can scale-out independently
- The storage and visualization are decoupled.
- Store once, visualize as required

# Open source projects Monitoring / Alerting

## Collector Agent



## Time Series Database



elasticsearch

## Alerting



Kapacitor

Elastalert

## Visualization

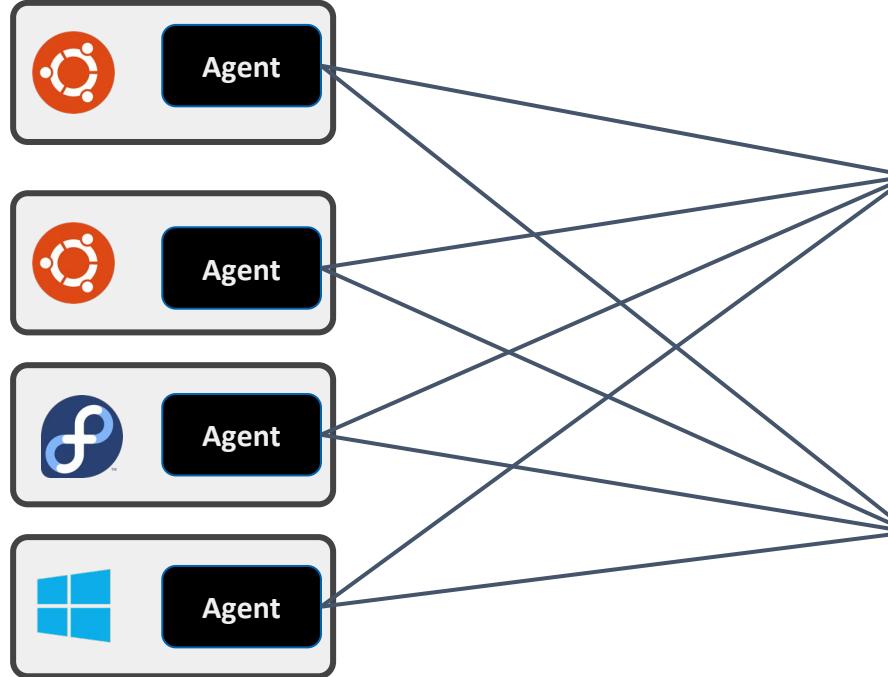


# Telegraf - The Swiss Army Knife

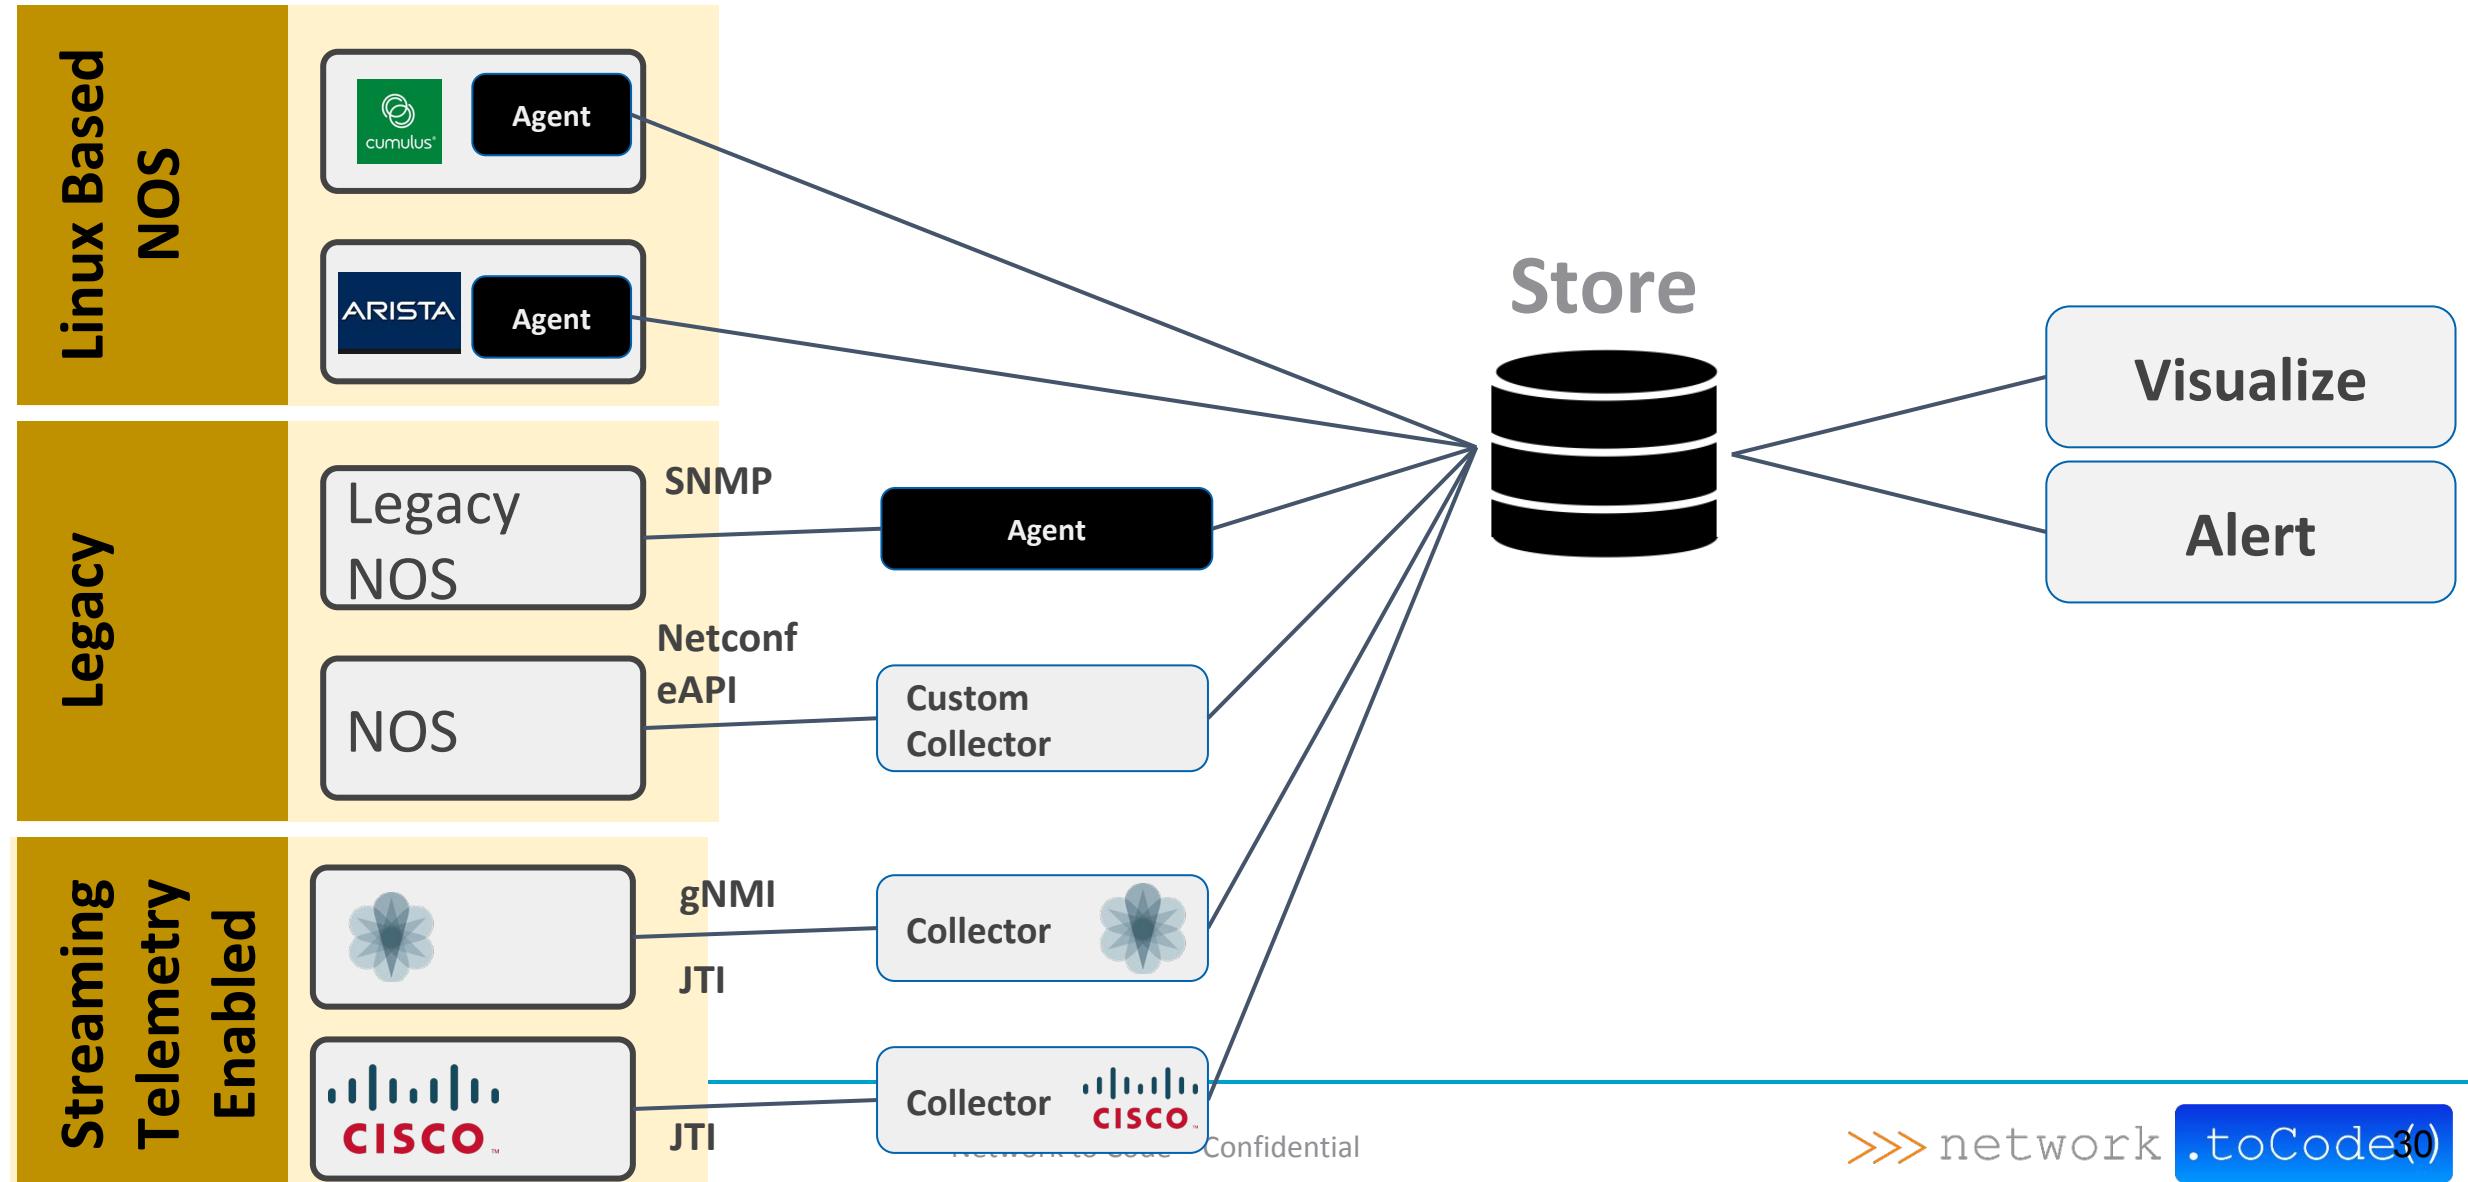


- Plugins driven agent / Extensible
- Support out of the box
  - Over 190 Input Plugins
  - Most Databases (output)
  - Data manipulation
- SNMP Input Plugin
- Juniper / Cisco / OpenConfig

# Cloud Based Solutions



# Reuse the same components for network devices



3

# Introduction to Time Series Database

# Modern Time Series Database



- **New generation of database optimized for Time serie data**
- **Started around 2013,  
Mainstream since 2016**
- **Powerful query engine**
- **Decorrelate storage and visualization**

# Introduction to Modern TSDB

`interface_output_bytes{device="spine1",interface="et-0/0/4"} 4569765412`

measurement name

What is it ?

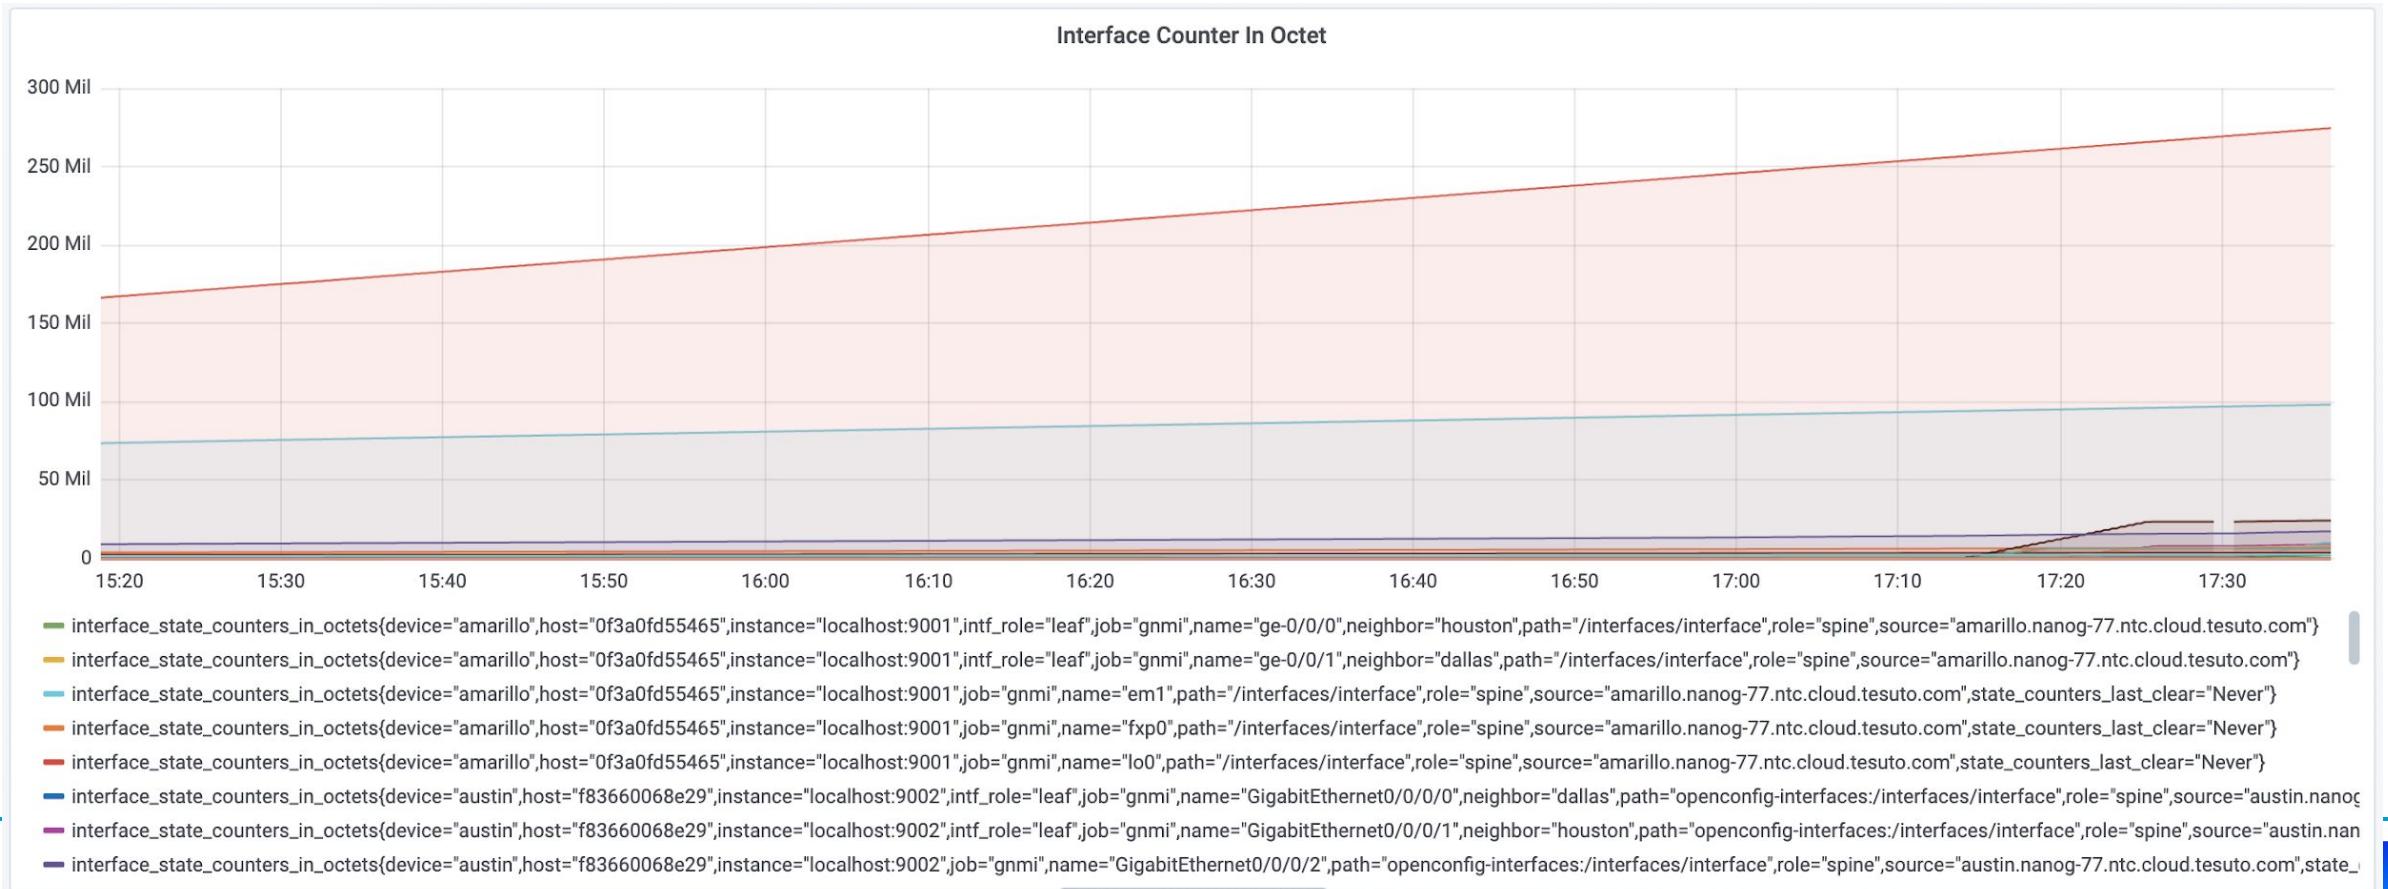
Tags/Labels

Contextual information

Value

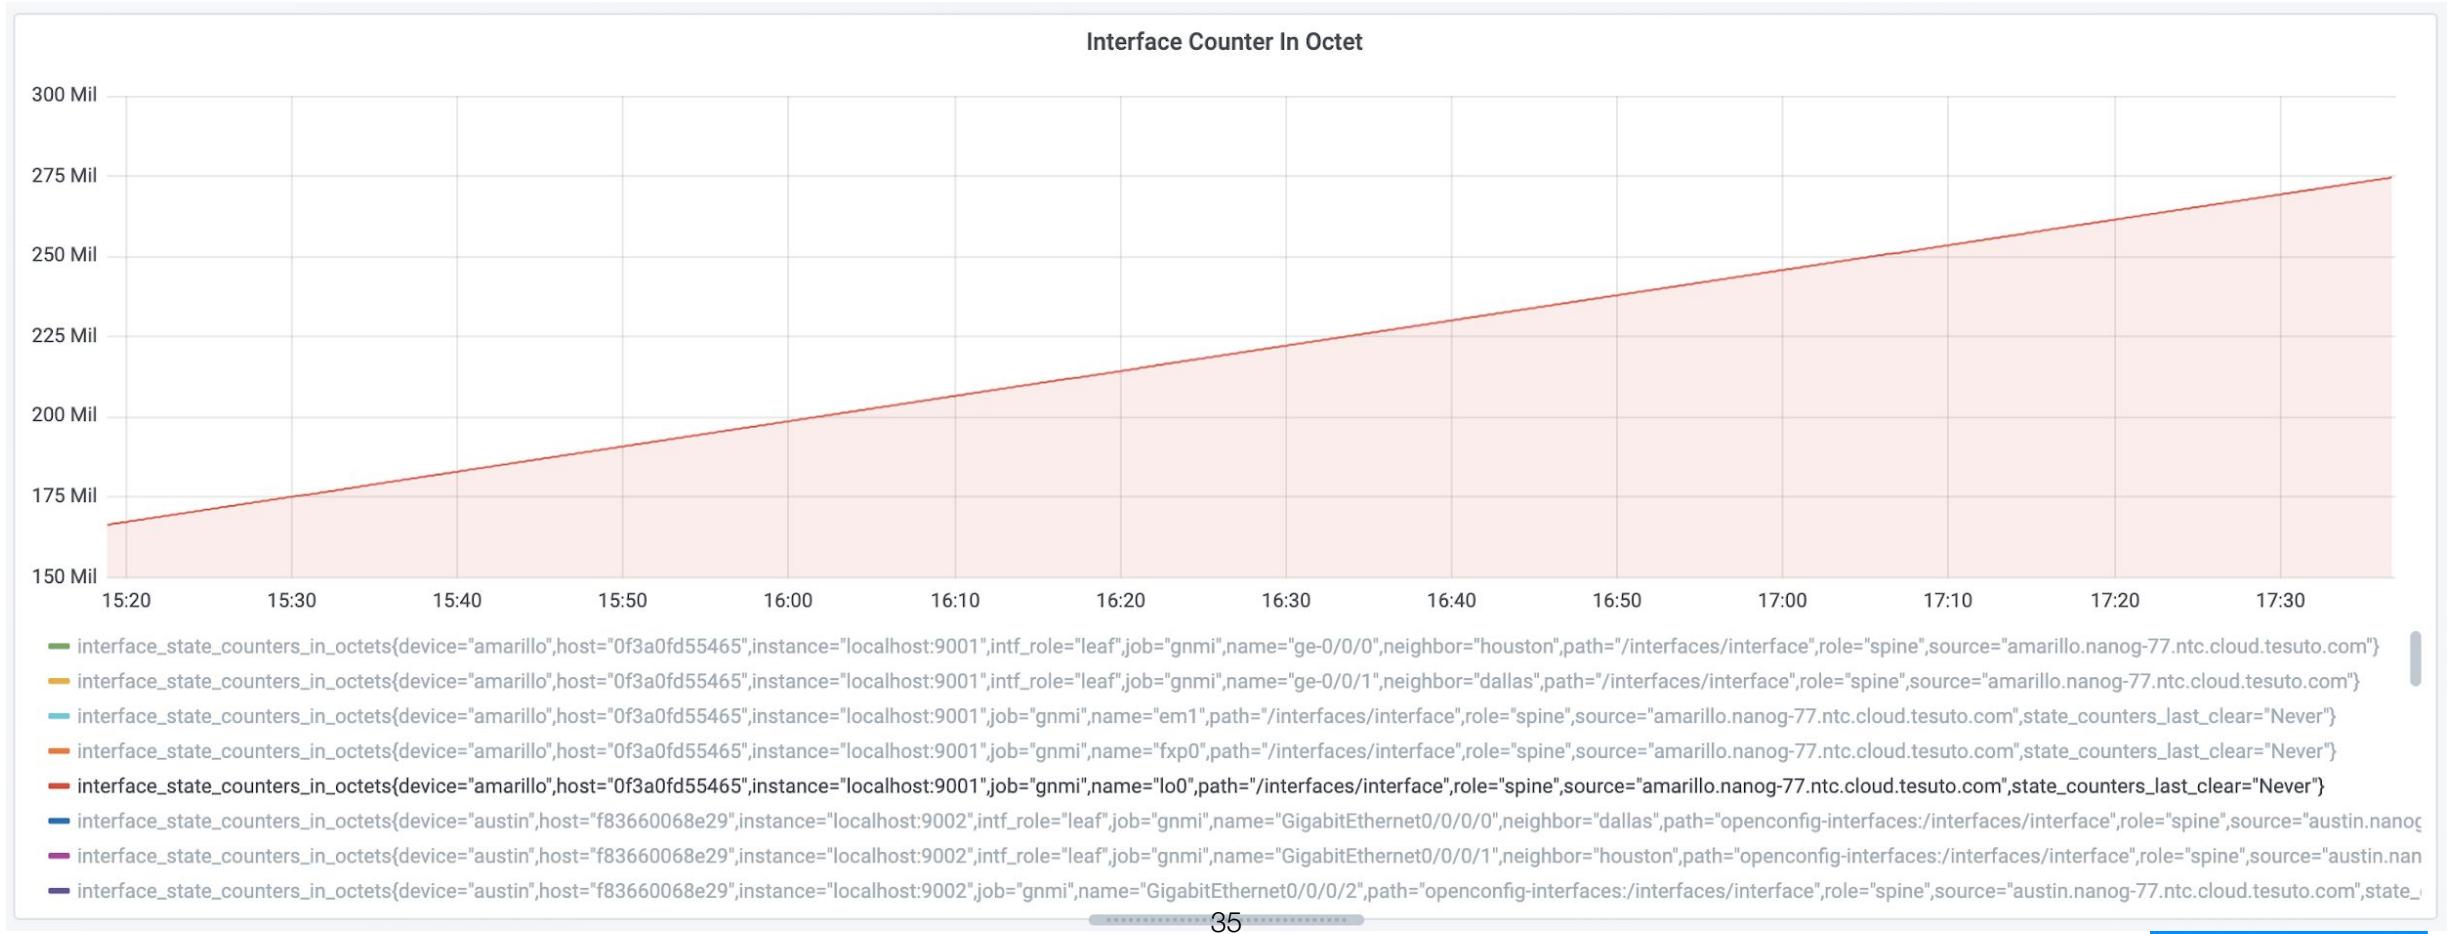
# Introduction to Modern TSDB

## interface\_state\_counters\_in\_octets



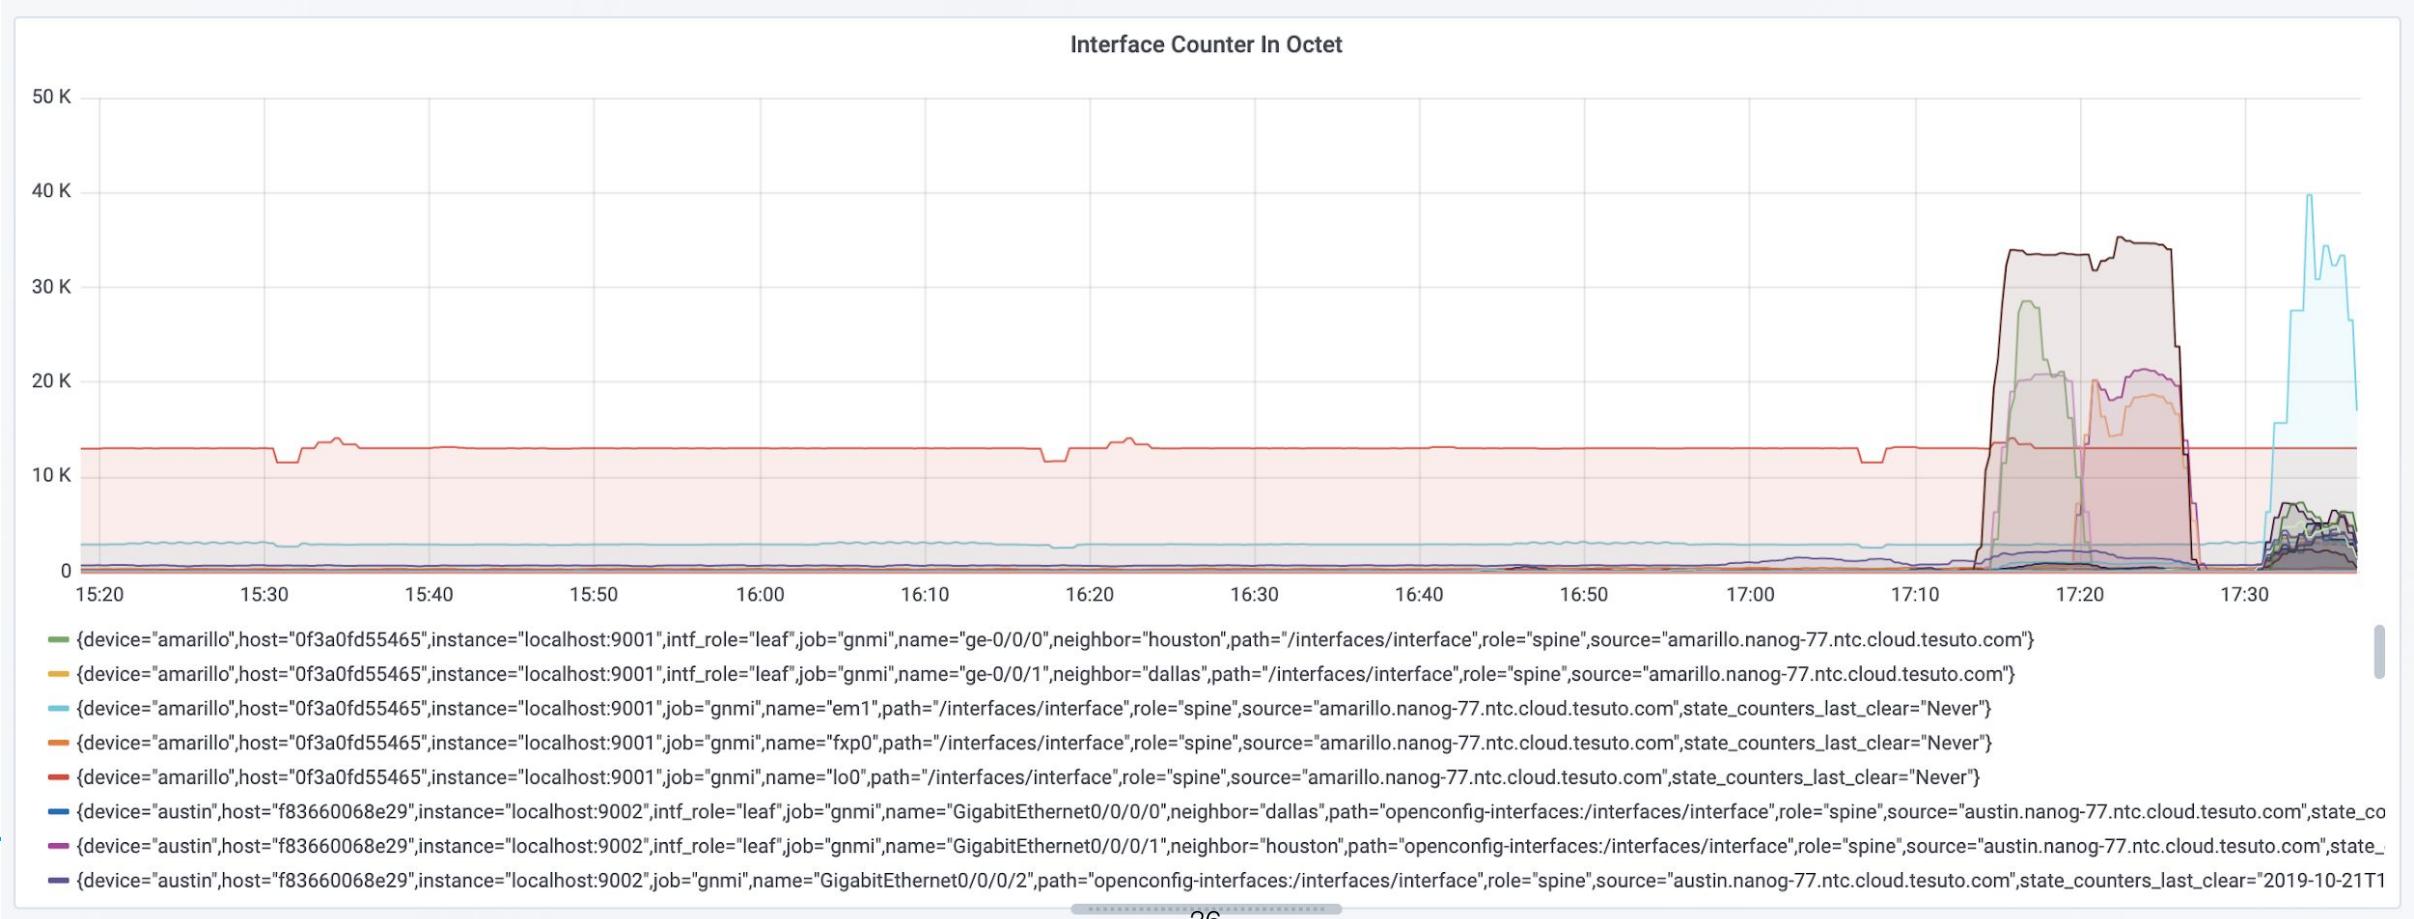
# Introduction to Modern TSDB

## interface\_state\_counters\_in\_octets



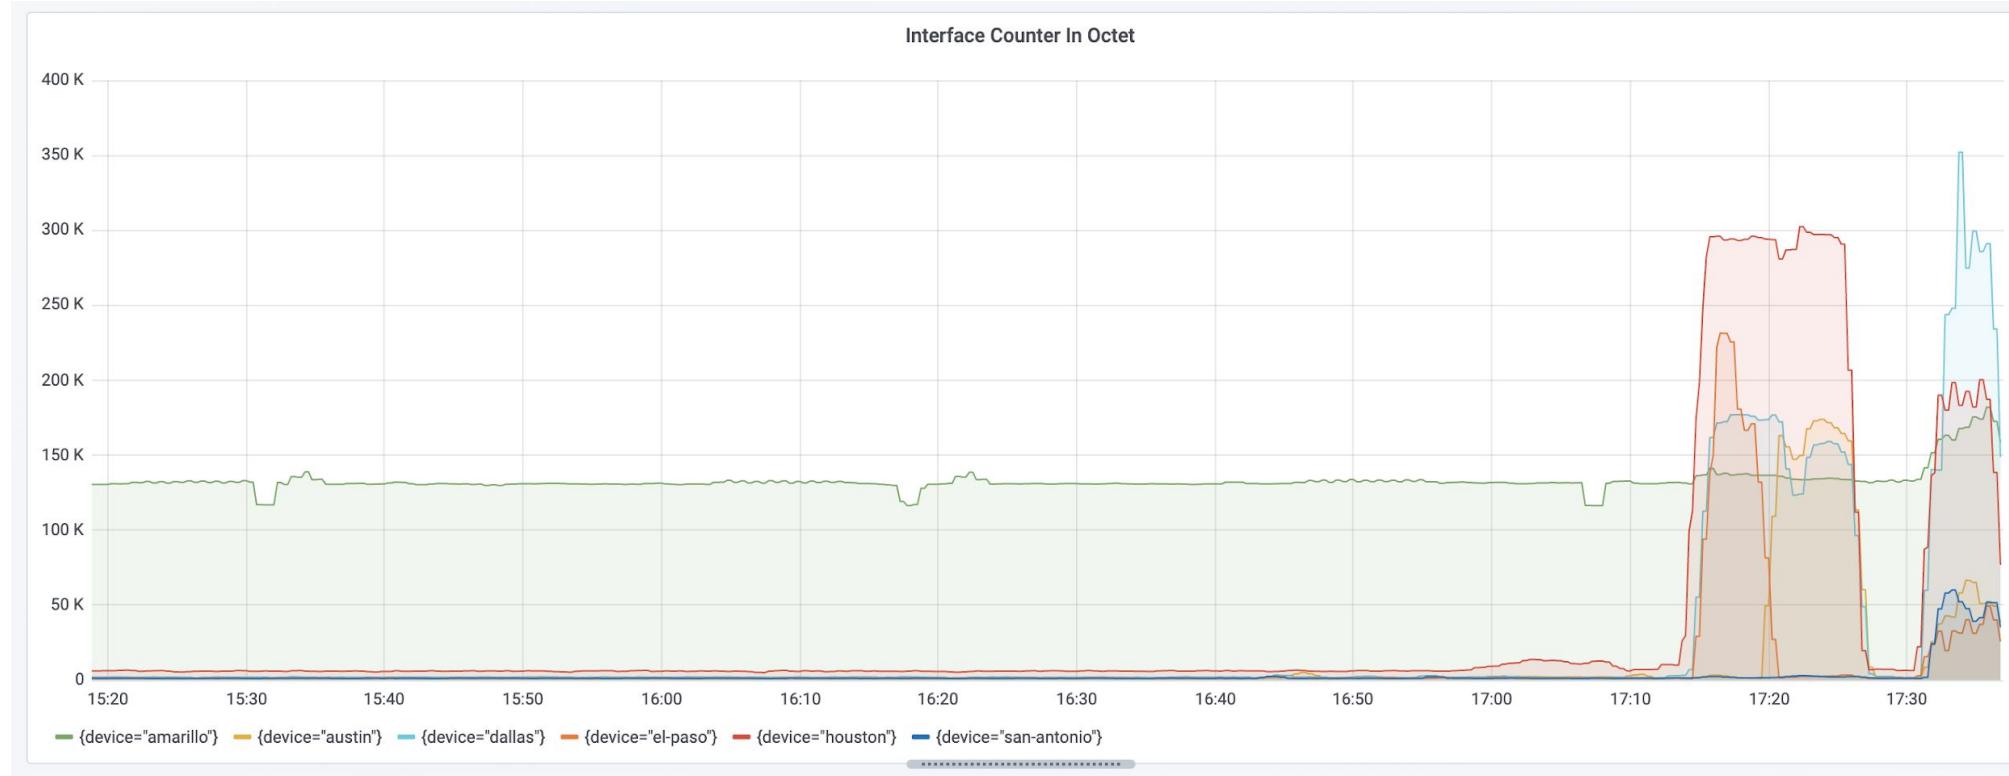
# Introduction to Modern TSDB

```
rate(interface_state_counters_in_octets[2m])*8
```



# Introduction to Modern TSDB

```
sum by (device)(  
    rate(interface_state_counters_in_octets[2m])  
)
```



# Introduction to Modern TSDB

```
interface_state_counters_in_octets{device="spine1",interface="et-0/0/4"} 4569765412
```

```
interface_state_counters_in_octets{  
    device="spine1", interface="et-0/0/4",  
    role="leaf",  
    site="fra1",  
    provider="level3",  
    intf_role="uplink"  
} 4569765412
```

# 4

## Building a modern monitoring solution from scratch

# How to Get Started

1. Select a Stack
2. Build an inventory
3. Collect Data
4. Add Metadata
5. Build Dashboards
6. Create Alerts

# Telemetry Stack



Grafana

**Dashboard**



Thanos



Prometheus

**Database**



*telegraf*

**Collector**

# Collect Data with Telegraf

snmp

gnmi

Execd

**Input Plugins**  
190+ available

regex

enum

**Processors Plugins**  
25+ available

prometheus\_client

influxdb

**Output Plugin**  
40+ available

# Few Telegraf Plugins

Name	Plugin Type	Description
<b>prometheus_client</b>	Output	Expose all data in prometheus format
<b>snmp</b>	Input	Collect information via SNMP
<b>gnmi</b>	Input	Collect information via gNMI / Telemetry Streaming
<b>Exec / execd</b>	Input	Execute third party script in any language (Python + Netmiko)
<b>dns_query</b>	Input	Measure DNS query response time
<b>net_response</b>	Input	Measure TCP response time
<b>http_response</b>	Input	Measure HTTP query response time
<b>rename</b>	Processor	Rename field and or tags
<b>regex</b>	Processor	Add / Drop / Rename field and tag based on Regex
<b>enum</b>	Processor	Convert text value into integer (UP > 1, DOWN > 0)

# Telegraf - SNMP

```
[[inputs.snmp]]
agents = ["device1"]
version = 2
community = "<SNMP Community>"           ← Base config : Credential, interval
interval = "60s"
timeout = "10s"
retries = 3

[[inputs.snmp.field]]
name = "hostname"
oid = "RFC1213-MIB::sysName.0"             ← Convert sysName.0 as tag hostname
is_tag = true

[[inputs.snmp.table]]
name = "interface"
inherit_tags = [ "hostname" ]                ← Collect IF-MIB::ifTable Table
oid = "IF-MIB::ifTable"

[[inputs.snmp.table.field]]
name = "name"
oid = "IF-MIB::ifDescr"                      ← Convert ifDescr as Tag name
is_tag = true
```

# Telegraf - gNMI

```
[inputs.gnmi]
addresses = ["device1"]
username = "<USERNAME>"
password = "<PASSWORD>"
```

Base config : Credential, interval

```
[[inputs.gnmi.subscription]
origin = "openconfig-interfaces"
path = "/interfaces/interface"
subscription_mode = "sample"
sample_interval = "60s"
```

Collect Interfaces/interface counters

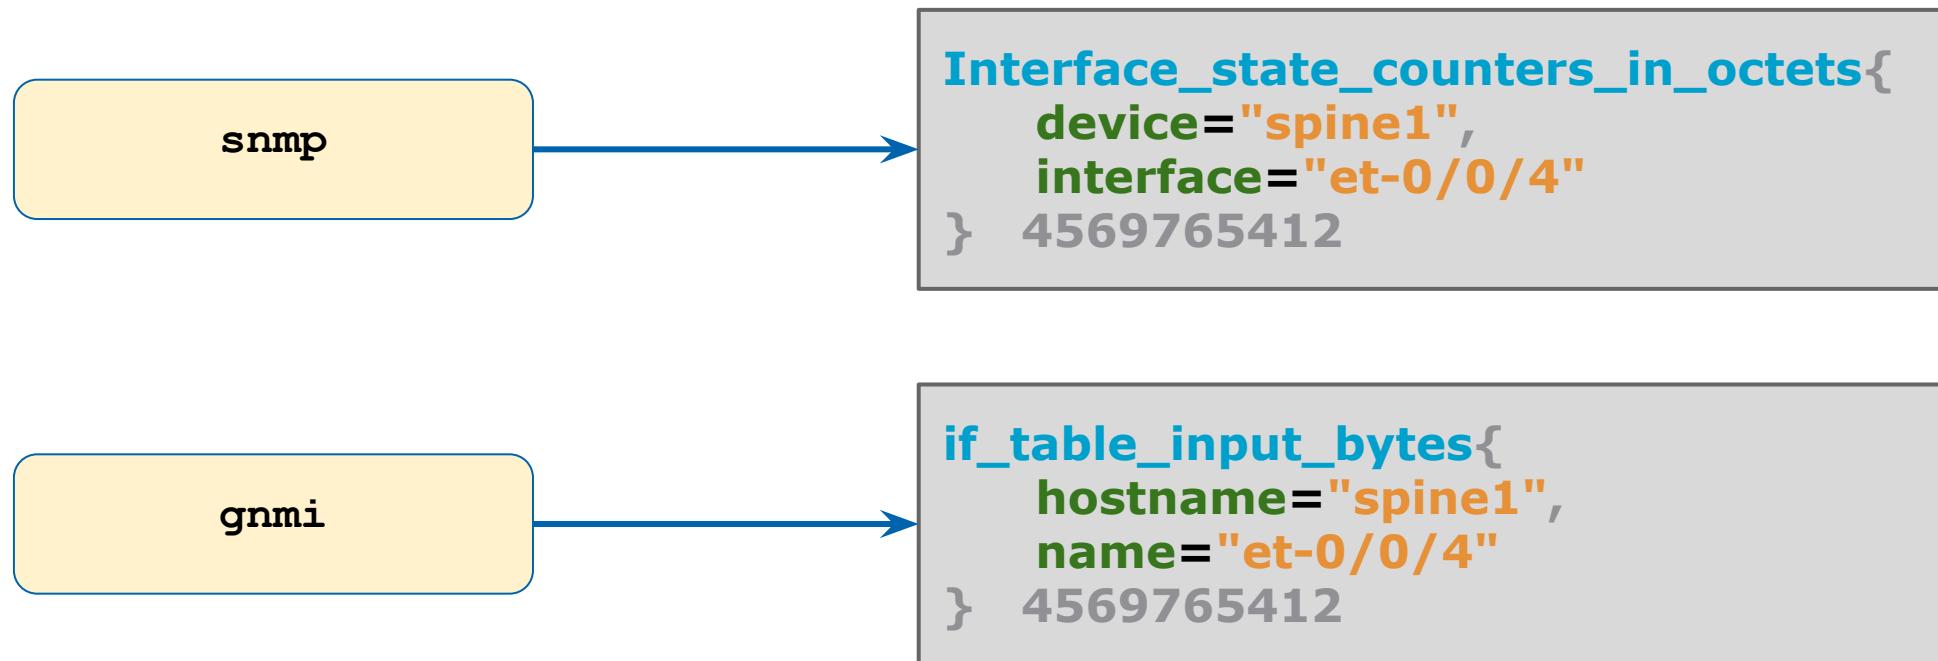
```
[[inputs.gnmi.subscription]]
name = "bgp_neighbor"
origin = "openconfig-network-instance"
path = "/network-instances/network-instance/protocols/protocol/bgp/neighbors/neighbor/state"
subscription_mode = "sample"
sample_interval = "10s"
```

Collect bgp neighbor state/counters

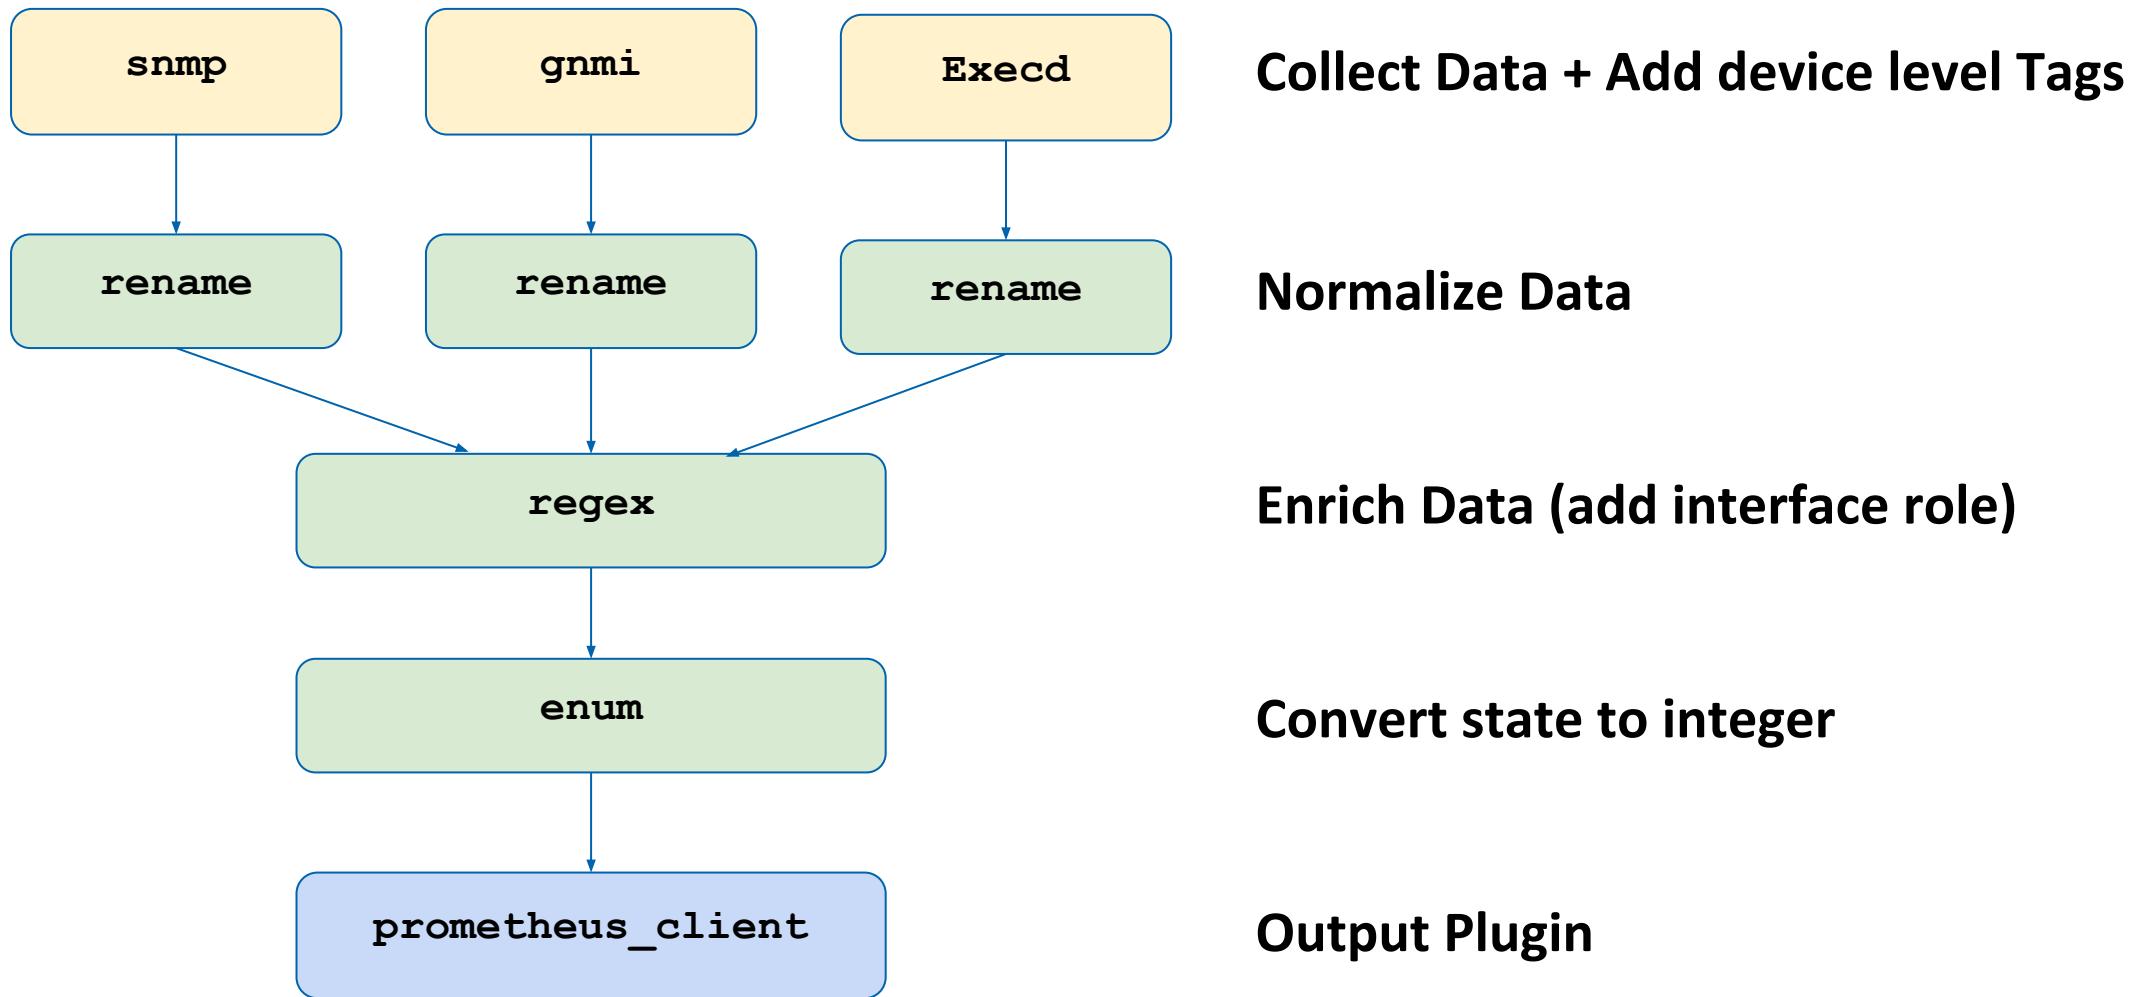
# Telegraf - execd

```
[[inputs.execd]]
interval = "60s"
command = [
    "python", "/source/my_collector.py", "--host", "device1",
]
signal = "SIGHUP"
restart_delay = "10s"
data_format = "influx"
```

# Challenge with Multiple Input Plugins



# Data Normalization with Telegraf - single device



# Telegraf - Add metadata

```
[ [processors.regex]]  
  
[processors.regex.tagpass]  
host = device1  
  
[ [processors.regex.tags]]  
key = "name"  
pattern = "^ge-0\\/0\\/0$"  
replacement = "management"  
result_key = "interfacerole"  
  
[ [processors.regex.tags]]  
key = "name"  
pattern = "^ge-0\\/0\\/1$"  
replacement = "data"  
result_key = "interfacerole"
```

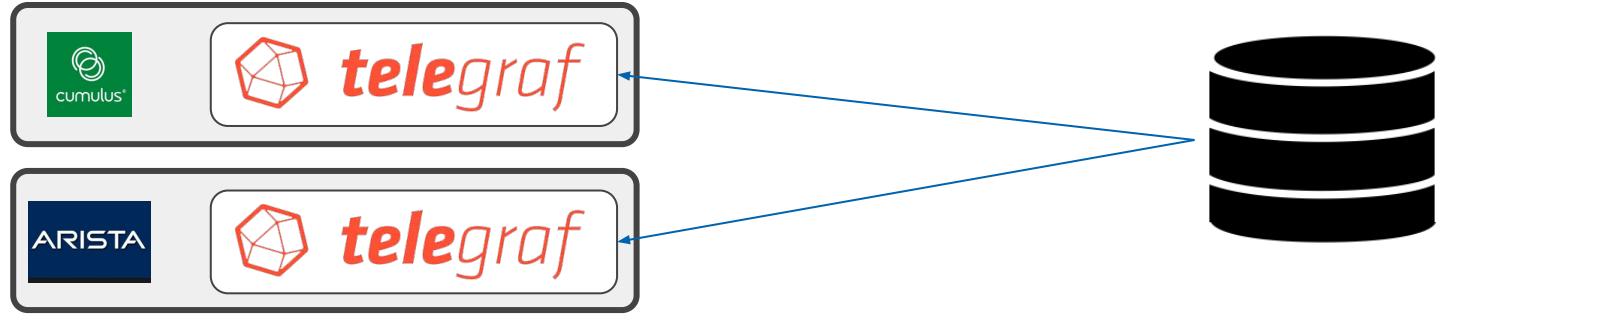
Apply these rules only for measurement with a tag host=device1

If a tag name=ge-0/0/0 is present, add a tag interfacerole=management

If a tag name=ge-0/0/1 is present, add a tag interfacerole=data

# Telegraf Deployment Methods

On-Box



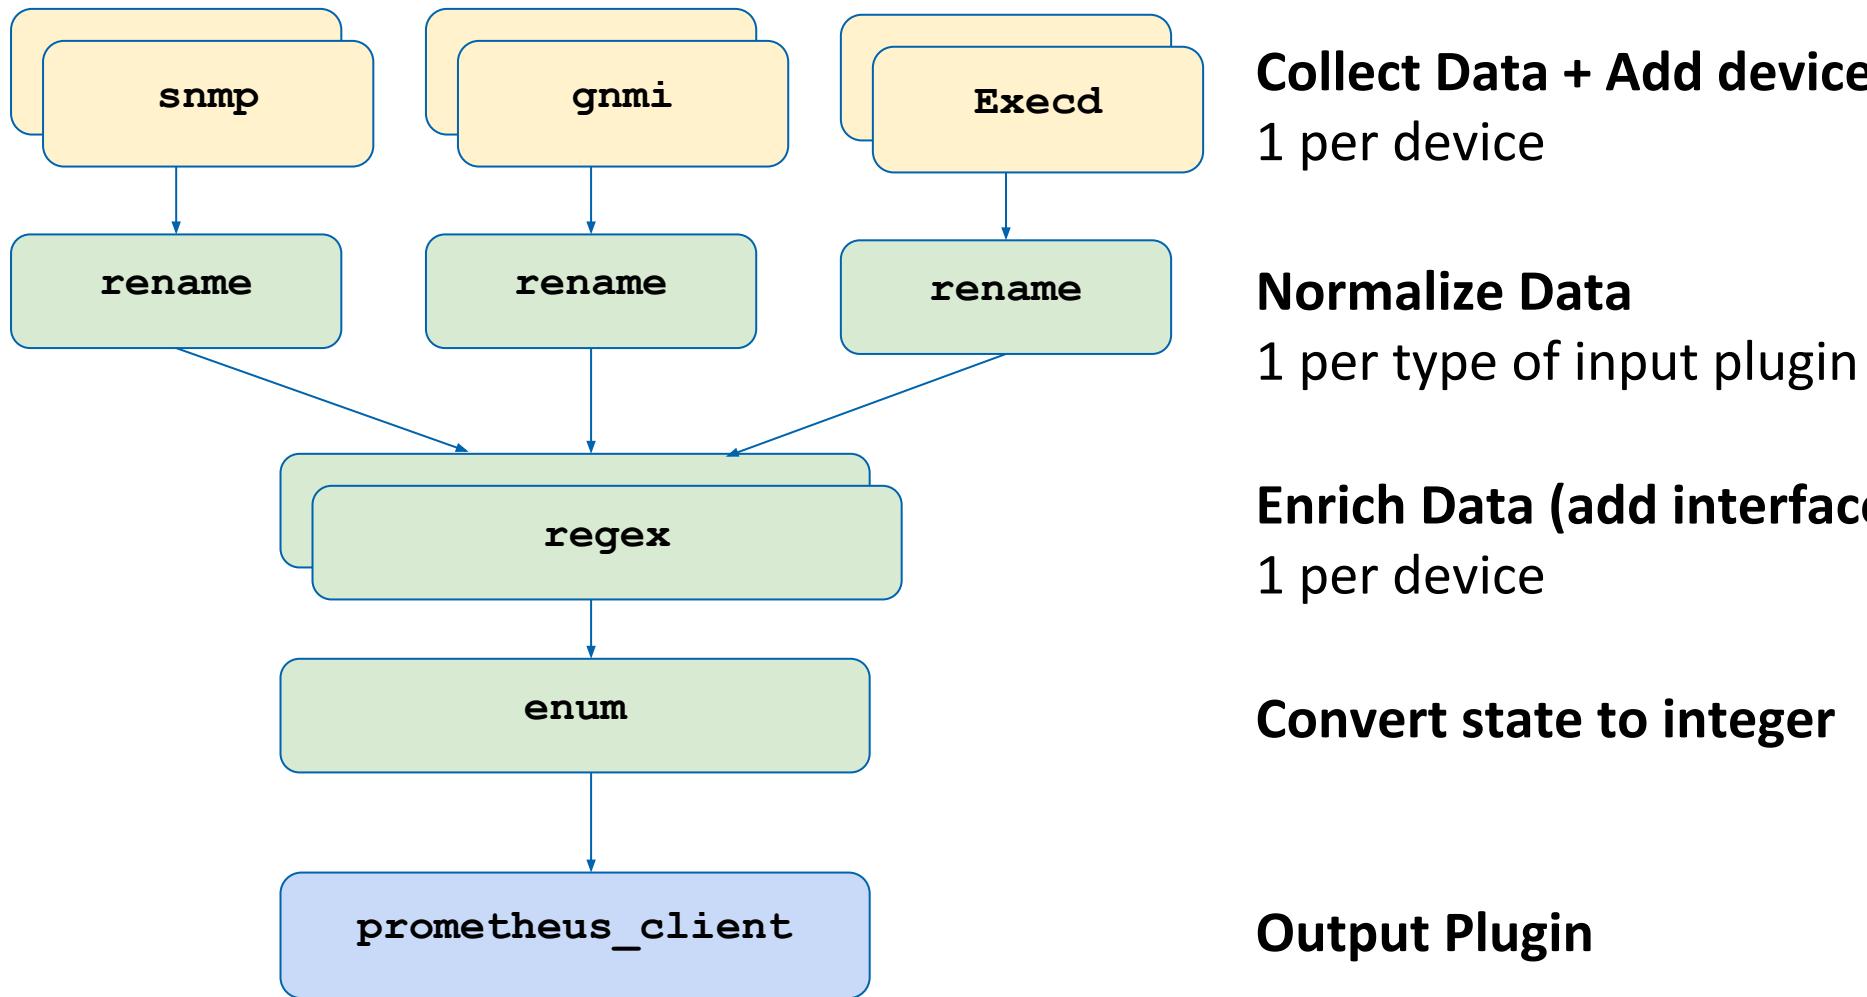
Off-Box - Dedicated



Off-Box - Shared



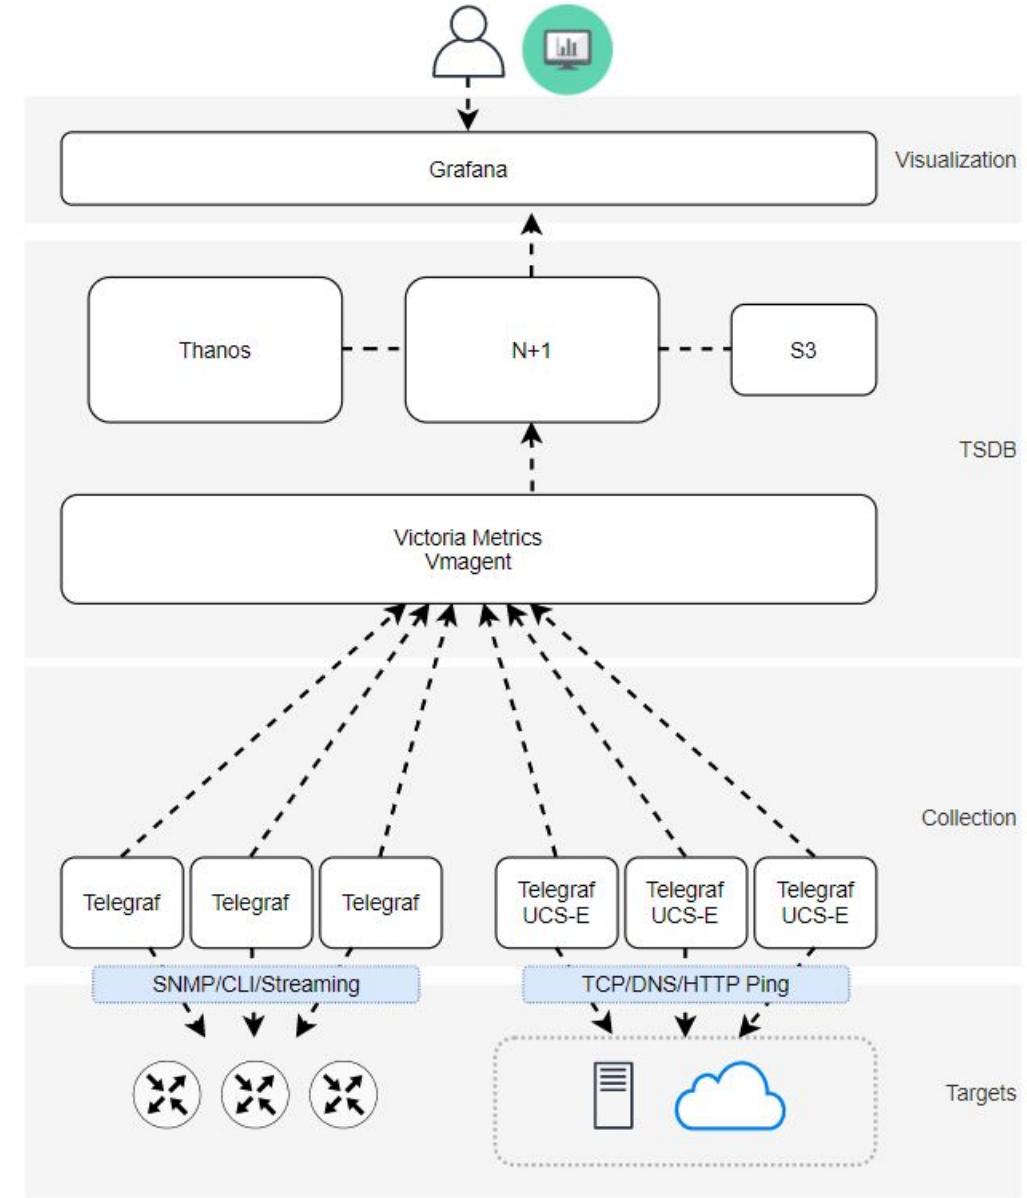
# Data Normalization with Telegraf - Multi-devices



# Target Architecture for Baxter and Lab Environment

# Target Architecture

- Target architecture is split into 4 main sections:
  - Visualization
  - TSDB
  - Collection
  - Targets
- Thanos. Prometheus solution, that provides:
  - long term storage via AWS S3.
  - HA.
- Visualization, TSDB and Vmagent components hosted on AWS ECS.
- Prometheus will be used to:
  - collect metrics from Thanos.
  - prevent circular dependency.



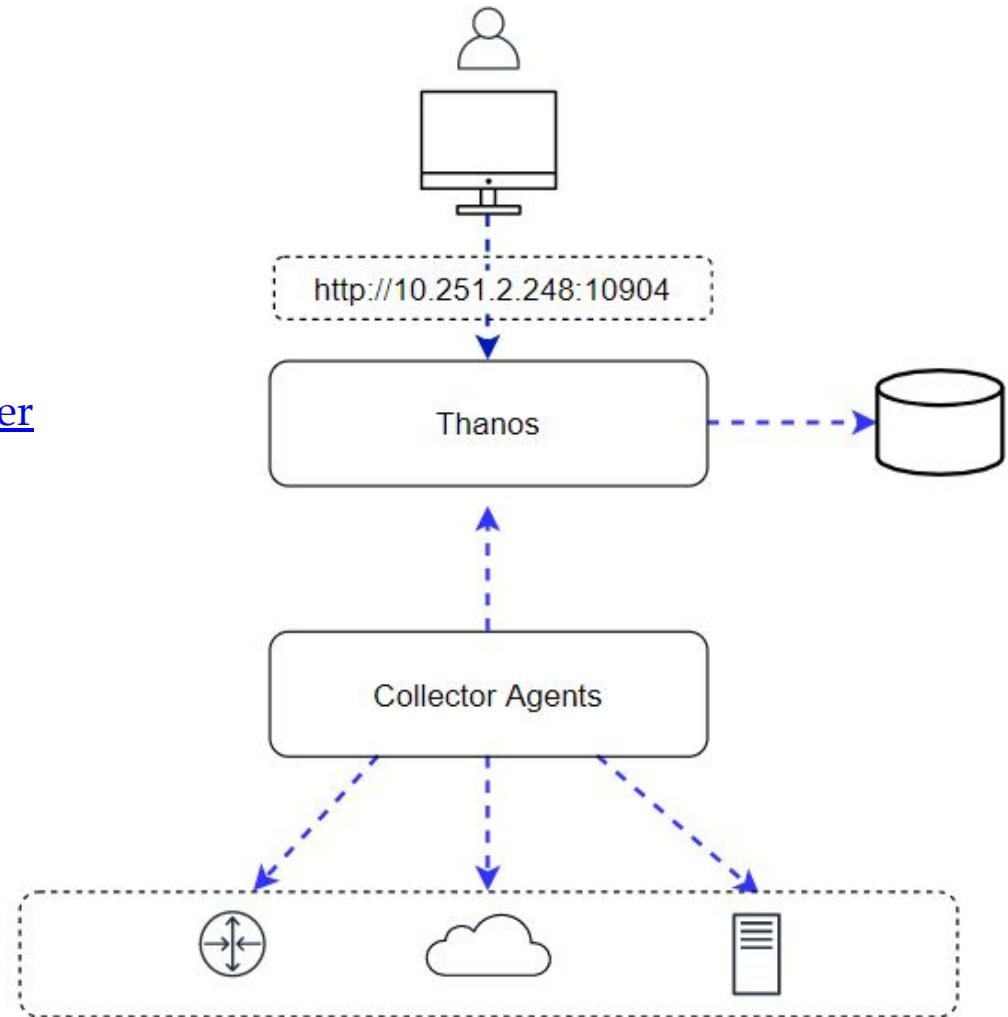
# Lab Environment

## Dashboards

- Thanos: <http://10.251.2.248:10904>
- Grafana: [network-telemetry-dev.aws.baxter.com](http://network-telemetry-dev.aws.baxter.com)

## Labs

- Repo: <https://github.com/networktocode/enablement-program-baxter>
- Path: workshops/telemetry-deepdive/prometheus\_promql/



# Data Available

Data	Collector	Metrics
<b>Linux Server monitoring</b>	Telegraf running on Linux	cpu_* disk_*
<b>DNS Response Time</b>	Telegraf running on UCSE	dns_query_query_time_ms
<b>Interface Counters</b>	Telegraf using SNMP for few routers in Europe	interface_in_* interface_out_*
<b>Netbox monitoring</b>	metric endpoint on Netbox	django_* netbox_*
<b>Zscaler Monitoring</b>	Custom collector written for Zscaler	zscaler_*
<b>ASA monitoring</b>	Telegraf agent executing python script via CLI	asa_*

# Prometheus & PromQL

# Overview

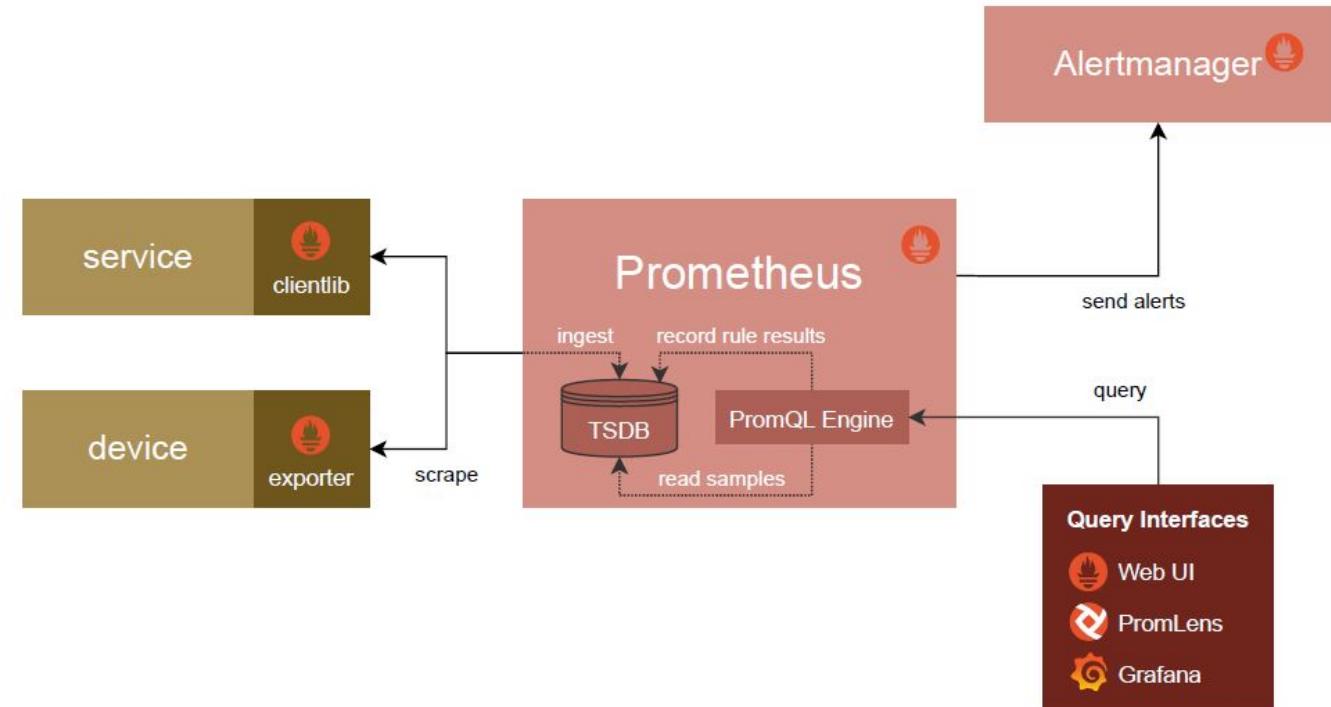
# What is Prometheus & PromQL?

## Prometheus

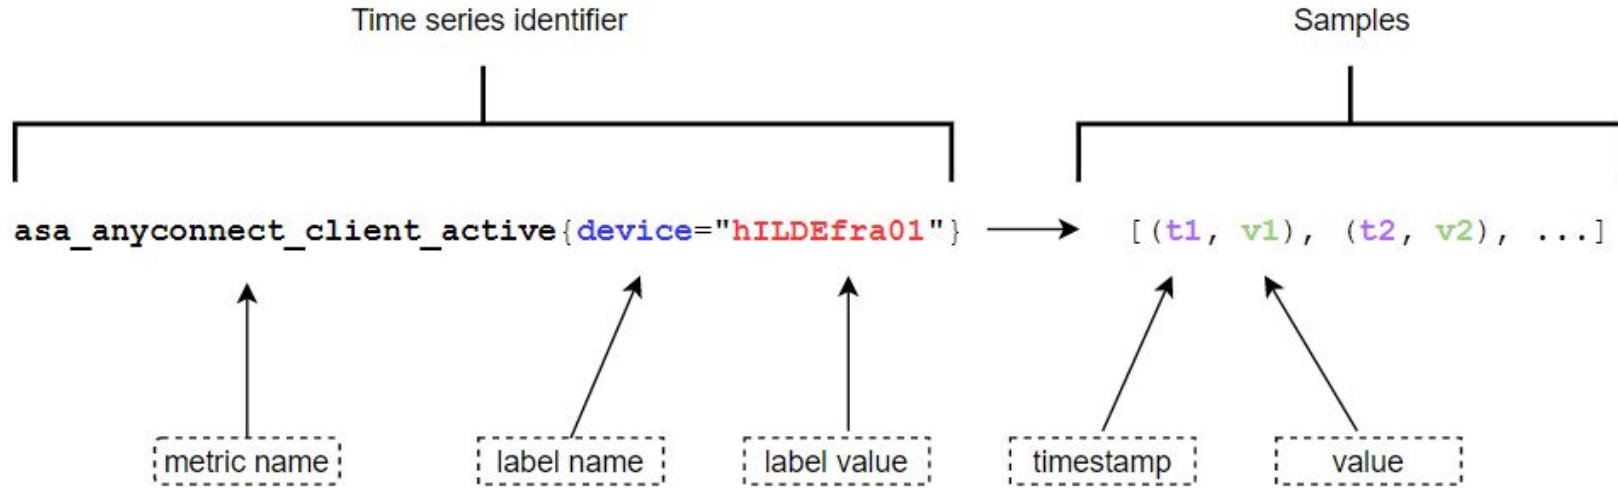
- Open-source systems monitoring and alerting toolkit.
- Includes:
  - TSDB (Time Series DB)
  - Time series data identified by metric name and key/value pairs.
  - PromQL query language.

## PromQL (Prometheus Query Language)

- Query language of the Prometheus monitoring system.
- Allows you to select, and aggregate time series data in realtime.
- Read-only.
- Use-cases: dashboards, alerting and ad-hoc querying.



# Time Series Data Model



- Prometheus stores all data as time series - streams of time stamped values.
  - Every time series is **uniquely identified** by its **metric name** and **optional key-value pairs** called labels.
  - The **metric name** specifies the general feature of a system that is measured (e.g. `asa_anyconnect_client_active` - the total number of active anyconnect client sessions)
  - **Labels** help to differentiate subdimensions (such as `site="china"`).
  - **Time series**: streams of time stamped values.

# Metric Types - Counters and Gauges

**Counters** - A value that only goes up.

Example: an interface counter.

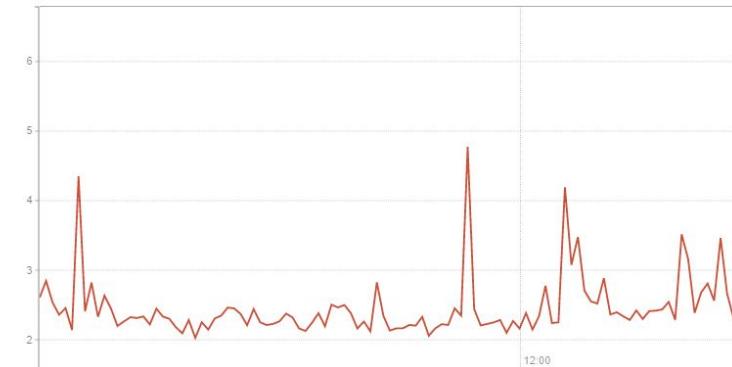
```
# TYPE net_packets_sent counter
net_packets_sent{host="deftlnetops200",interface="eth0"}
```



**Gauges** - A value that go down as well as up.

Example: CPU usage.

```
# TYPE disk_total gauge
cpu_usage_guest{cpu="cpu-total",host="deftlnetops200"}
```



# Metric Types - Histograms and Summaries

## Histograms

- Sampled values are put into buckets each as individual time series.
- Buckets such as “requests under 5ms”, “requests under 10ms” etc.
- Quantiles can then be calculated via `histogram_quantile()`.
- Histograms typically used for request latencies or response times.

```
# TYPE prometheus_http_response_size_bytes histogram
prometheus_http_response_size_bytes_bucket{handler="/" ,le="100"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="1000"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="10000"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="100000"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="1e+06"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="1e+07"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="1e+08"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="1e+09"} 12
prometheus_http_response_size_bytes_bucket{handler="/" ,le="+Inf"} 12
prometheus_http_response_size_bytes_sum{handler="/" } 348
prometheus_http_response_size_bytes_count{handler="/" } 12
```

## Summaries

- Predecessor to Histograms.
- Essentially the same as histograms but the the quantiles are calculated on the client side.
- Client side quantiles are sent as time-series.

```
# TYPE prometheus_target_interval_length_seconds summary
prometheus_target_interval_length_seconds{interval="10s",quantile="0.01"} 9.999700111
prometheus_target_interval_length_seconds{interval="10s",quantile="0.05"} 9.999953889
prometheus_target_interval_length_seconds{interval="10s",quantile="0.5"} 10.000012499
prometheus_target_interval_length_seconds{interval="10s",quantile="0.9"} 10.000066784
prometheus_target_interval_length_seconds{interval="10s",quantile="0.99"} 10.000355232
```

# Queries and Filtering

# Thanos/Prometheus UI

Thanos Graph Stores Status ▾ Help New UI

Enable query history

asa\_anyconnect\_client\_active ← PromQL query //

Load time: 64ms  
Resolution: 14s  
Total time series: 12

**Execute** - insert metric at cursor -  deduplication  partial response

Graph Console

◀ Moment ▶

Element	Value
asa_anyconnect_client_active{device="hILDEfra02",environment="ip-10-251-2-248.aws.baxter.com",instance="localhost:9471",job="asa_vpn",role="asa_vpn",site="frankfurt",tenant_id="default-tenant"}	161
asa_anyconnect_client_active{device="hILUSder02",environment="ip-10-251-2-248.aws.baxter.com",instance="localhost:9477",job="asa_vpn",role="asa_vpn",site="deerfield",tenant_id="default-tenant"}	1
asa_anyconnect_client_active{device="hILUSchi02",environment="ip-10-251-2-248.aws.baxter.com",instance="localhost:9474",job="asa_vpn",role="asa_vpn",site="elk_grove_ch3",tenant_id="default-tenant"}	252
asa_anyconnect_client_active{device="hILSGsie03",environment="ip-10-251-2-248.aws.baxter.com",instance="localhost:9481",job="asa_vpn",role="asa_vpn",site="singapore_sg1",tenant_id="default-tenant"}	41
asa_anyconnect_client_active{device="hILSGsie01",environment="ip-10-251-2-248.aws.baxter.com",instance="localhost:9479",job="asa_vpn",role="asa_vpn",site="singapore_sg1",tenant_id="default-tenant"}	55

# Label Filtering

Allows you to match a label value, or to match label values against regular expressions.

- `=` Select labels that are **equal** to the provided string.
- `!=` Select labels that are **not equal** to the provided string.
- `=~` Select labels that **regex-match** the provided string.
- `!~` Select labels that do **not regex-match** the provided string.

```
# equal
interface_ifInOctets{device="zpa_frankfurt_fr6_1"}

# not equal
interface_ifInOctets{name!="lo"}

# regex-match
interface_ifInOctets{device=~"zpa_aws_frankfurt_[3-4] "}

# not regex-match
interface_ifInOctets{device!~"zpa_aws_frankfurt_[3-4] "}
```

Lab001

# Offsets

Offset Modifiers allow you to change the time offset for the returned timeseries.

Allows you to view samples in the past.

Example time durations:

- 5h
- 1h30m
- 5m
- 10s

```
# admin interface status 48 hours ago
interface_ifAdminStatus{device="zpa_frankfurt_fr6_1",
name="eth0"} offset 48h
```

# Instant vs Range Vector Selection

- Some **PromQL functions** require a range of time-series for each series over time, instead of a single value.
- Time-series metrics can be expressed in 2 ways:
  - **Instant vector** - time series identifier (metric and labels) returns one value.
  - **Range vector** - time series identifier (metric and labels) returns values for a given period.
- These are also known as **Expression Types**.

```
// Instant Vector  
asa_anyconnect_client_active  
55  
  
// Range Vector  
asa_anyconnect_client_active[5m]  
56 @1605557193.498  
56 @1605557223.498  
56 @1605557253.498  
56 @1605557283.498  
56 @1605557313.498  
56 @1605557343.498  
55 @1605557373.498  
55 @1605557403.498  
55 @1605557433.498  
55 @1605557463.498
```

Lab003

# Prometheus Functions

# Prometheus Functions

Prometheus functions typically applies to the values returned from a series. Such as:

- Instant vector
- Range vector

The screenshot shows a Prometheus query interface with the following details:

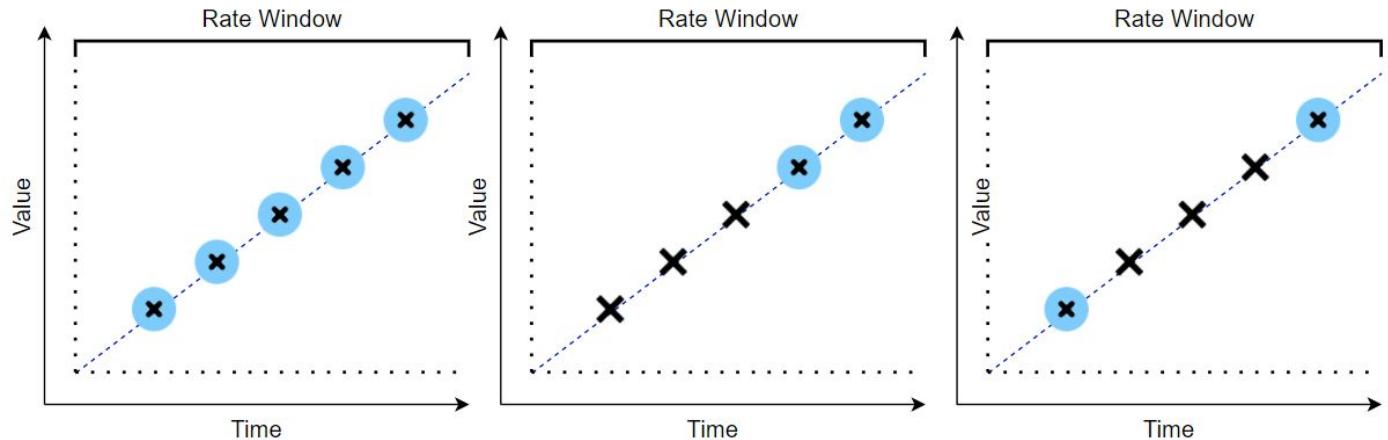
- Query:** asa\_anyconnect\_client\_active[2m]
- Buttons:** Execute, - insert metric at cursor -, deduplication, partial response
- Graph/Console:** Graph is selected.
- Time Range:** Moment
- Table Results:**

Element	Value
asa_anyconnect_client_active{device="h1DEfra02",environment="ip-10-251-2-248.aws.baxter.com",instance="localhost:9471",job="asa_vpn",role="asa_vpn",site="frankfurt",tenant_id="default-tenant"}	155 @1605776320.517 155 @1605776350.517 155 @1605776380.517 155 @1605776410.517

A red box highlights the last four rows of the table results. A dashed box labeled "Lab004" is located in the bottom right corner of the results area.

# rate(), irate(), increase()

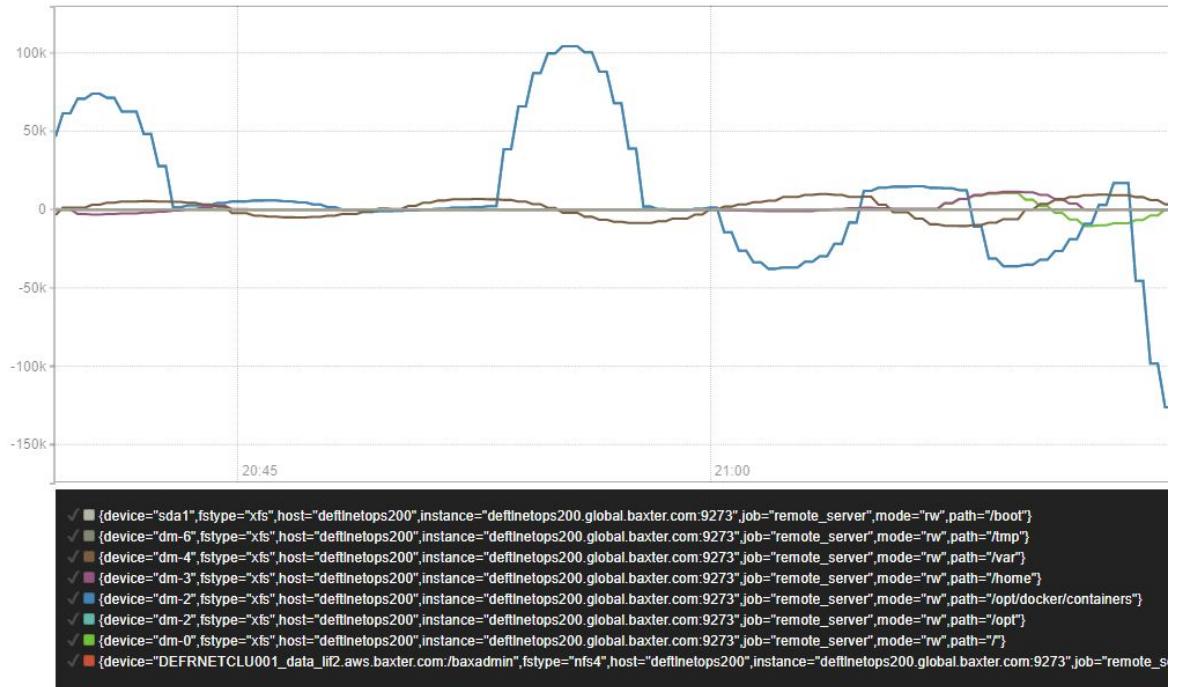
- Counter metrics will only ever increase.
- This results in graphs that just go up and up.
- Prometheus provides 3 functions to graph increase:
  - **rate()** - calculates the **per-second increase** of a counter as averaged over a time range. (range vector)
  - **irate()** - calculates the **per-second increase** between last 2 points over a time range (range vector).
  - **increase()** - the **total increase** over a time range (range vector).
- Requirements: at least 2 samples.
- Recommendations: range window size at least 4x scrape interval.



```
// rate() - provides smoother results. Better for alerting  
rate(interface_ifHCOutUcastPkts{...} [5m])  
  
// irate() - useful for high resolution graphing  
irate(interface_ifHCOutUcastPkts{...} [5m])  
  
// increase()  
increase(interface_ifHCOutUcastPkts{...} [5m])
```

# deriv(), delta()

- Used with gauge based metrics (value that can go up or down).
- Prometheus provides 2 functions that allow us to graph these changes:
  - **deriv()** - calculates the per-second derivative of a set of series over a specified time range (range vector).
  - **delta()** - calculates the difference between the first and last value over a specified time range (range vector).



```
# deriv()
deriv(disk_used{host="deftlnetops200"} [5m])

# delta()
delta(disk_used{host="deftlnetops200"} [5m])
```

Lab005

# <aggregation>\_over\_time()

- Takes a range vector (series of metric values).
- Performs the necessary aggregation against the range vector series.
  - **avg\_over\_time(range-vector)**: the average value of all points in the specified interval.
  - **min\_over\_time(range-vector)**: the minimum value of all points in the specified interval.
  - **max\_over\_time(range-vector)**: the maximum value of all points in the specified interval.
  - **sum\_over\_time(range-vector)**: the sum of all values in the specified interval.
  - **count\_over\_time(range-vector)**: the count of all values in the specified interval.

```
> asa_anyconnect_client_active{device="hILDEFra01"} [5m]
58 @1605640136.309
57 @1605640166.309
62 @1605640196.309
62 @1605640226.309
62 @1605640256.309
57 @1605640286.309
57 @1605640316.309
43 @1605640346.309
53 @1605640376.309
57 @1605640406.309

> max_over_time(asa_anyconnect_client_active{device=
" hILDEFra01"} [5m])
62
```

# Operators

# Operators

- Binary operators
  - Arithmetic binary operators (+, -, \*, /)
  - Comparison binary operators (==, !=, >, <)
  - Logical/set binary operators (and, or, unless)
- Vector matching
  - Aggregation operators (sum, min, max)
  - One-to-one vector matches
  - Many-to-one and one-to-many vector matches

```
rate(interface_in_octets{hostname="rAWITgro01", intf_type="mpls"} [5m]) * 8  
/  
interface_speed{hostname="rAWITgro01", intf_type="mpls"}
```

# sum(), avg(), count()

**Aggregation Operators** aggregate the elements of a single instant vector. Examples include:

- sum() - Adds together the instant vectors returned from multiple samples.
- avg() - Calculates an average of the instant vectors returned from multiple samples.
- count() - Counts the number of elements in the vector.

! Not to be confused with <aggregation>\_over\_time(), which is a function.

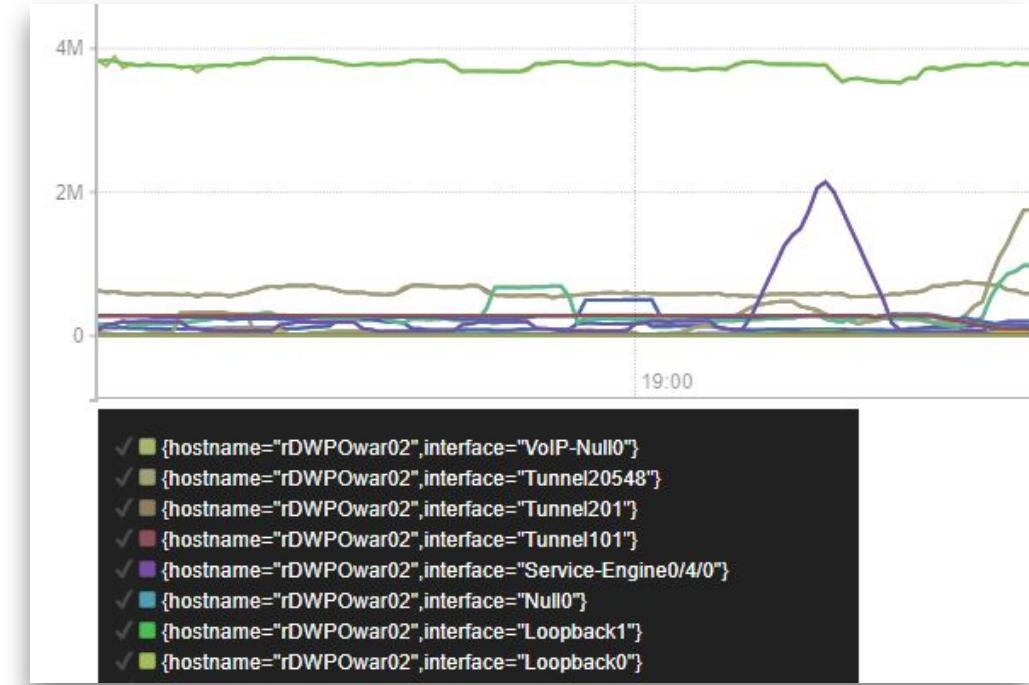
```
// sum()
sum(asa_anyconnect_client_active{site="frankfurt"})

// avg()
avg(asa_anyconnect_client_active{site="frankfurt"})

// count()
count(asa_anyconnect_client_active{site="frankfurt"})
```

# by(), without() Clauses

- Clauses can be used to preserve distinct dimensions when running the operator.
- They are combined with the operator.
- The 2 clauses available are:
  - by() - includes the labels listed by the clause.
  - without() - exclude the labels listed by the clause.



```
// by()
sum by (hostname,interface) (rate(interface_in_octets{} [5m]))

// without()
sum without (job,agent_host) (rate(interface_in_octets{} [5m]))
```

Lab008

# Advanced Topics

# Quantile over Time

- Quantile over time is an <aggregation>\_over\_time() function.
- Allows you to calculate the quantile from a range vector.
- Takes the percentile, along with a range vector.
- For example:
  - Show me the top 10% slowest DNS query times from a 5 minute time range?

```
// quantile_over_time(scalar, range-vector)
quantile_over_time(0.90, dns_query_query_time_ms{site="lessines", server=~"10.100.13[0-1].61"} [5m]
```

Lab009

# Standard Deviation

There are 2 ways to calculate the standard deviation of your time series values.

- `stddev_over_time(range-vector)` (function)
  - a standard deviation is calculated over the values within a range vector
- `stddev()` (operator)
  - performed over each sample across multiple vectors.

	Value
:dns-	23.372168 @1605737410 22.274869 @1605737420 21.971859 @1605737430 22.86833 @1605737440 20.539724 @1605737450
c-dns-b.google.com",site="braines-	19.126187 @1605737410 18.809977 @1605737420 19.567863 @1605737430 20.57715 @1605737440 19.59154 @1605737450
nes",tenant_id="default-	23.673075
nes",tenant_id="default-	19.897601
	21.629957

```
// results in the stddev for dns ms for each device in the site lessines
stddev_over_time(dns_query_query_time_ms{site="lessines", sensor="one.one.one.one"} [5m])

// results in the stddev for dns ms across all devices in the site in the site lessines
stddev(dns_query_query_time_ms{site="lessines", sensor="one.one.one.one"})
```

Lab009

# Function - predict\_linear()

- Used with gauge based metrics (value that can go up or down).
- Predicts the value of a gauge in the future based on a time range (range vector).
- Requires the number of seconds from now.

```
# predict_linear(v range-vector, t scalar)

sum by (hostname,interface) (predict_linear(interface_in_octets{}[5m], 2592000)) * 8

sum by (device) (predict_linear(interface_ifInUcastPkts{}[5m], 25920))
```